
Final In-Class Exam (Solution Ideas)

(2:35 PM – 3:50 PM : 75 Minutes)

- This exam will account for either 15% or 30% of your overall grade depending on your relative performance in the midterm and the final. The higher of the two scores (midterm and final) will be worth 30% of your grade, and the lower one 15%.
- There are three (3) questions, worth 75 points in total. Please answer all of them in the spaces provided.
- There are 16 pages including three (3) blank pages and two (2) pages of appendices. Please use the blank pages if you need additional space for your answers.
- The exam is *open slides* and *open notes*.

GOOD LUCK!

Question	Pages	Score	Maximum
1. Parallel Prefix Sum	2–5		25
2. ϵ -Approximate Frequency	7–9		30
3. Matrix Rotation	11–13		20
Total			75

NAME: _____

QUESTION 1. [25 Points] Parallel Prefix Sum. Given a sequence of n elements $\langle x_1, x_2, \dots, x_n \rangle$ drawn from a set S with a binary associative operator \oplus (e.g., addition, multiplication, maximum, matrix product, union, etc.), the *prefix sum* problem asks one to compute a sequence of n partial sums $\langle s_1, s_2, \dots, s_n \rangle$ such that $s_i = x_1 \oplus x_2 \oplus \dots \oplus x_i$ for $1 \leq i \leq n$. In lecture 26 we studied a parallel prefix sum algorithm with $\Theta(n)$ work and $\Theta(\log^2 n)$ span¹.

In this problem we will analyze another parallel prefix sum algorithm given in Figure 1.

```

ALT-PREFIX-SUM(  $\langle x_1, x_2, \dots, x_n \rangle, \oplus$  )
(Input is a sequence of  $n$  elements  $\langle x_1, x_2, \dots, x_n \rangle$  and a binary associative operator  $\oplus$ . Output is a sequence
 $\langle s_1, s_2, \dots, s_n \rangle$  with  $s_i = x_1 \oplus x_2 \oplus \dots \oplus x_i$ , for  $1 \leq i \leq n$ . We assume  $n = 2^k$  for some integer  $k \geq 0$ .)

1. if  $n = 1$  then
2.      $s_1 \leftarrow x_1$                                      {the prefix sum of a single element is the element itself}
3. else
4.     spawn  $\langle s_1, s_2, \dots, s_{\frac{n}{2}} \rangle \leftarrow \text{ALT-PREFIX-SUM} \left( \langle x_1, x_2, \dots, x_{\frac{n}{2}} \rangle, \oplus \right)$ 
                                                {sets  $s_i = x_1 \oplus x_2 \oplus \dots \oplus x_i$  for  $1 \leq i \leq \frac{n}{2}$ }
5.      $\langle s_{\frac{n}{2}+1}, s_{\frac{n}{2}+2}, \dots, s_n \rangle \leftarrow \text{ALT-PREFIX-SUM} \left( \langle x_{\frac{n}{2}+1}, x_{\frac{n}{2}+2}, \dots, x_n \rangle, \oplus \right)$ 
                                                {sets  $s_{\frac{n}{2}+i} = x_{\frac{n}{2}+1} \oplus x_{\frac{n}{2}+2} \oplus \dots \oplus x_{\frac{n}{2}+i}$  for  $1 \leq i \leq \frac{n}{2}$ }
6.     sync
7.     parallel for  $i \leftarrow 1$  to  $\frac{n}{2}$  do
8.          $s_{\frac{n}{2}+i} \leftarrow s_{\frac{n}{2}} \oplus s_{\frac{n}{2}+i}$ 
                                                {extends  $s_{\frac{n}{2}+i} = x_{\frac{n}{2}+1} \oplus x_{\frac{n}{2}+2} \oplus \dots \oplus x_{\frac{n}{2}+i}$ 
to  $s_{\frac{n}{2}+i} = s_{\frac{n}{2}} \oplus x_{\frac{n}{2}+1} \oplus x_{\frac{n}{2}+2} \oplus \dots \oplus x_{\frac{n}{2}+i} = x_1 \oplus x_2 \oplus \dots \oplus x_{\frac{n}{2}+i}$ }
9. return  $\langle s_1, s_2, \dots, s_n \rangle$ 

```

Figure 1: An alternate parallel prefix sum algorithm.

¹assuming the span of a *parallel for* loop with n iterations to be $\mathcal{O}(\log n + k)$, where k is the maximum span of a single iteration

1(a) [**7 Points**] Write down a recurrence relation describing the work done (i.e., T_1) by ALT-PREFIX-SUM, and solve it.

Solution. Let $T_1(n)$ be the work done by the algorithm for a sequence of n elements.

Clearly, the algorithm performs $\Theta(1)$ work (in line 2) when $n = 1$. Otherwise (i.e., if $n > 1$), it makes two recursive calls to itself (in lines 4 and 5) each for a sequence of length $\frac{n}{2}$, and then executes a **for** loop (in lines 6 and 7) that iterates $\frac{n}{2}$ times. Hence, T_1 can be described using the following recurrence relation.

$$T_1(n) = \begin{cases} \Theta(1) & \text{if } n = 1, \\ 2T_1\left(\frac{n}{2}\right) + \Theta(n) & \text{otherwise.} \end{cases}$$

Using Master Theorem (case 2), we obtain: $T_1(n) = \Theta(n \log n)$.

- 1(b) [**7 Points**] Write down a recurrence relation describing the span (i.e., T_∞) of ALT-PREFIX-SUM, and solve it.

Solution. Let $T_\infty(n)$ be the span of the algorithm for an input of size n .

Clearly, the algorithm has $\Theta(1)$ span (in line 2) for $n = 1$. Otherwise (i.e., if $n > 1$), it makes two parallel recursive calls to itself (in lines 4 and 5) each for a sequence of length $\frac{n}{2}$, and then executes a ***parallel for*** loop (in lines 6 and 7) that iterates $\frac{n}{2}$ times and thus has $\Theta(\log n)$ span. Hence, T_∞ can be described using the following recurrence relation.

$$T_\infty(n) = \begin{cases} \Theta(1) & \text{if } n = 1, \\ T_\infty\left(\frac{n}{2}\right) + \Theta(\log n) & \text{otherwise.} \end{cases}$$

Using Master Theorem (case 2), we obtain: $T_\infty(n) = \Theta(\log^2 n)$.

1(c) [**6 Points**] Find the parallel running time (i.e., T_p) and parallelism of ALT-PREFIX-SUM.

Solution.

Parallel running time with p processing elements on an input of size n (using Graham / Brent theorem): $T_p(n) = \Theta\left(\frac{T_1(n)}{p} + T_\infty(n)\right) = \Theta\left(\left(\frac{n}{p} + \log n\right) \log n\right)$.

Parallelism for an input of size n : $P(n) = \frac{T_1(n)}{T_\infty(n)} = \Theta\left(\frac{n}{\log n}\right)$.

1(d) [**5 Points**] Is ALT-PREFIX-SUM work-optimal? Why or why not?

Solution. ALT-PREFIX-SUM is work-optimal provided $pT_p(n) = \Theta(T_S(n))$, where $T_S(n)$ is the work performed by an optimal serial prefix sum algorithm.

We know that for an input of length n , any prefix sum algorithm must perform $\Omega(n)$ work (since it must read the input at least once). We also know that one can compute the prefix sums in $\mathcal{O}(n)$ time simply by scanning the input sequence once from left to right once and always keeping track of the running sum. Hence, $T_S(n) = \Theta(n)$.

On the other hand, for ALT-PREFIX-SUM, $pT_p = \Theta(n \log n + p \log^2 n) = \omega(T_S(n))$.

Hence, the algorithm is not work-optimal.

Use this page if you need additional space for your answers.

QUESTION 2. [30 Points] ϵ -Approximate Frequency. Let $A[1 : n]$ be an array of length n containing both positive and negative numbers. Let m be the number of positive numbers in A , and let $p = \frac{m}{n}$. We are interested in estimating the value of m fast. Clearly, one can find the exact value of m in $\Theta(n)$ time simply by scanning A once and counting the number of positive numbers.

For any $\epsilon \in (0, p]$, we say that \hat{m} is an ϵ -approximation² of m provided $m - \epsilon n < \hat{m} < m + \epsilon n$.

APPROX-FREQ($A[1 : n]$, ϵ)

(Inputs are an array $A[1 : n]$ of n numbers, and a floating point parameter $\epsilon \in (0, 1]$. This routine chooses a sample of size $\lceil \frac{6}{\epsilon^2} \ln n \rceil$ from A uniformly at random (with replacement), and uses that sample to estimate the number of entries of A that are positive.)

1. $s \leftarrow \lceil \frac{6}{\epsilon^2} \ln n \rceil$ {size of the sample}
2. $c \leftarrow 0$ {a counter that keeps track of the frequency of v in the chosen sample}
3. **for** $i \leftarrow 1$ **to** s **do** {sample s items (with replacement) from A }
4. $j \leftarrow \text{RANDOM}(1, n)$ {choose an integer uniformly at random from $[1, n]$ }
5. **if** $A[j] > 0$ **then** $c \leftarrow c + 1$ {choose $A[j]$ as the next sample from A }
6. **return** $\frac{c}{s} \times n$ {return the estimate}

Figure 2: Estimate the number of entries of $A[1 : n]$ that are positive.

This problem asks you to show that the function APPROX-FREQ given in Figure 2 which runs in $\Theta(\frac{1}{\epsilon^2} \ln n)$ worst-case time returns an ϵ -approximation of m w.h.p. in n . While analyzing the algorithm we will drop the ceiling in line 1 for simplicity, i.e., we will assume that $s = \frac{6}{\epsilon^2} \ln n$.

2(a) [5 Points] Let μ be the expected value of c right after the loop in lines 3–5 completes execution. Show that $\mu = (\frac{6}{\epsilon^2}) (\frac{m}{n}) \ln n$.

Solution. For $1 \leq i \leq s$, let X_i be a 0-1 random variable defined as follows.

$$X_i = \begin{cases} 1 & \text{if iteration } i \text{ samples a positive number,} \\ 0 & \text{otherwise.} \end{cases}$$

Clearly, $\Pr[X_i = 1] = \frac{m}{n} = p$, and $E[X_i] = 1 \times \Pr[X_i = 1] + 0 \times \Pr[X_i = 0] = p$.

Then $\mu = \sum_{i=1}^s E[X_i] = \sum_{i=1}^s p = sp = (\frac{6}{\epsilon^2}) (\frac{m}{n}) \ln n$.

²for simplicity, we have used ' $<$ ' instead of ' \leq ' in the definition of ϵ -approximation

2(b) [**12 Points**] Let \hat{c} be the exact value of c right after the loop in lines 3–5 completes execution. Prove that for $0 < \epsilon < p$ and $\delta = \frac{\epsilon}{p}$,

$$\Pr[\hat{c} \leq (1 - \delta)\mu] < \frac{1}{n^3} \quad \text{and} \quad \Pr[\hat{c} \geq (1 + \delta)\mu] < \frac{1}{n^2}.$$

Solution. Given $0 < \epsilon < p \Rightarrow 0 < \frac{\epsilon}{p} < 1 \Rightarrow 0 < \delta < 1$.

Observe from part 1(a) that $\hat{c} = \sum_{i=1}^s X_i$, where each X_i is a Poisson random variable.

We also have from part 1(a), $\mu = \left(\frac{6}{\epsilon^2}\right)\left(\frac{m}{n}\right) \ln n = \left(\frac{6}{\epsilon^2}\right)p \ln n$.

Hence, we can use the following Chernoff bound for the lower tail:

$$\begin{aligned} \Pr[\hat{c} \leq (1 - \delta)\mu] &\leq e^{-\frac{\mu\delta^2}{2}} \\ &= e^{-\left(\frac{6}{\epsilon^2}\right)p \ln n \times \left(\frac{\epsilon}{p}\right)^2 \times \frac{1}{2}} \\ &= e^{-\frac{3}{p} \ln n} \\ &= n^{-\frac{3}{p}} \\ &< n^{-3} && \left\{ \because p < 1 \Rightarrow \frac{1}{p} > 1 \Rightarrow -\frac{3}{p} < -3 \Rightarrow n^{-\frac{3}{p}} < n^{-3} \right\} \\ &= \frac{1}{n^3} \end{aligned}$$

For the upper tail we can use the following Chernoff bound:

$$\begin{aligned} \Pr[\hat{c} \geq (1 + \delta)\mu] &\leq e^{-\frac{\mu\delta^2}{3}} \\ &= e^{-\left(\frac{6}{\epsilon^2}\right)p \ln n \times \left(\frac{\epsilon}{p}\right)^2 \times \frac{1}{3}} \\ &= e^{-\frac{2}{p} \ln n} \\ &= n^{-\frac{2}{p}} \\ &< n^{-2} && \left\{ \because p < 1 \Rightarrow \frac{1}{p} > 1 \Rightarrow -\frac{2}{p} < -2 \Rightarrow n^{-\frac{2}{p}} < n^{-2} \right\} \\ &= \frac{1}{n^2} \end{aligned}$$

2(c) [**5 Points**] let \hat{m} be the estimate of m returned by APPROX-FREQ. Argue that for $0 < \epsilon < p$, the results from part 2(b) imply the following:

$$\Pr [\hat{m} \leq m - \epsilon n] < \frac{1}{n^3} \text{ and } \Pr [\hat{m} \geq m + \epsilon n] < \frac{1}{n^2}.$$

Solution. Observe from line 6 of APPROX-FREQ that $\hat{m} = \frac{\hat{c}}{s} \times n$. Then

$$\begin{aligned} \Pr [\hat{m} \leq m - \epsilon n] &= \Pr \left[\frac{\hat{c}}{s} \times n \leq (1 - \epsilon \times \frac{n}{m})m \right] \\ &= \Pr \left[\hat{c} \leq \left(1 - \epsilon \times \frac{n}{m}\right) s \left(\frac{m}{n}\right) \right] \\ &= \Pr \left[\hat{c} \leq \left(1 - \frac{\epsilon}{p}\right) sp \right] && \left\{ \because \frac{m}{n} = p \text{ (given)} \right\} \\ &= \Pr [\hat{c} \leq (1 - \delta)\mu] && \left\{ \because \frac{\epsilon}{p} = \delta \text{ (from 2(b)) and } sp = \mu \text{ (from 2(a))} \right\} \\ &< \frac{1}{n^3} && \left\{ \text{from part 2(b)} \right\} \end{aligned}$$

$$\begin{aligned} \text{and } \Pr [\hat{m} \geq m + \epsilon n] &= \Pr \left[\frac{\hat{c}}{s} \times n \geq (1 + \epsilon \times \frac{n}{m})m \right] \\ &= \Pr \left[\hat{c} \geq \left(1 + \epsilon \times \frac{n}{m}\right) s \left(\frac{m}{n}\right) \right] \\ &= \Pr \left[\hat{c} \geq \left(1 + \frac{\epsilon}{p}\right) sp \right] && \left\{ \because \frac{m}{n} = p \text{ (given)} \right\} \\ &= \Pr [\hat{c} \geq (1 + \delta)\mu] && \left\{ \because \frac{\epsilon}{p} = \delta \text{ (from 2(b)) and } sp = \mu \text{ (from 2(a))} \right\} \\ &< \frac{1}{n^2} && \left\{ \text{from part 2(b)} \right\} \end{aligned}$$

2(d) [**8 Points**] Use your results from part 2(c) to argue that for $0 < \epsilon < p$, APPROX-FREQ returns an ϵ -approximation of m w.h.p. in n .

Solution. We have:

$$\begin{aligned} \Pr [m - \epsilon n < \hat{m} < m + \epsilon n] &= \Pr [(\hat{m} > m - \epsilon n) \cap (\hat{m} < m + \epsilon n)] \\ &= \Pr [\hat{m} > m - \epsilon n] + \Pr [\hat{m} < m + \epsilon n] \\ &\quad - \Pr [(\hat{m} > m - \epsilon n) \cup (\hat{m} < m + \epsilon n)] \\ &= \Pr [\hat{m} > m - \epsilon n] + \Pr [\hat{m} < m + \epsilon n] - 1 \\ &\quad \left\{ \because \Pr [(\hat{m} > m - \epsilon n) \cup (\hat{m} < m + \epsilon n)] = 1 \text{ for } \epsilon > 0 \right\} \\ &= (1 - \Pr [\hat{m} \leq m - \epsilon n]) + (1 - \Pr [\hat{m} \geq m + \epsilon n]) - 1 \\ &= 1 - \Pr [\hat{m} \leq m - \epsilon n] - \Pr [\hat{m} \geq m + \epsilon n] \\ &> 1 - \frac{1}{n^3} - \frac{1}{n^2} \left\{ \text{from part 2(c)} \right\} \\ &\geq 1 - \frac{2}{n^2} \left\{ \because n^3 \geq n^2 \Rightarrow -\frac{1}{n^3} \geq -\frac{1}{n^2} \right\} \end{aligned}$$

Hence, APPROX-FREQ returns an ϵ -approximation of m w.h.p. in n .

Use this page if you need additional space for your answers.

QUESTION 3. [20 Points] Matrix Rotation. The *rotation* of an $n \times n$ matrix X is another $n \times n$ matrix X^R obtained by writing the i -th row of X as the $n - i + 1$ -th column of X^R for $1 \leq i \leq n$. An example is given below.

$$X = \begin{bmatrix} a_1 & a_2 & a_3 & a_4 \\ b_1 & b_2 & b_3 & b_4 \\ c_1 & c_2 & c_3 & c_4 \\ d_1 & d_2 & d_3 & d_4 \end{bmatrix} \Rightarrow X^R = \begin{bmatrix} d_1 & c_1 & b_1 & a_1 \\ d_2 & c_2 & b_2 & a_2 \\ d_3 & c_3 & b_3 & a_3 \\ d_4 & c_4 & b_4 & a_4 \end{bmatrix}$$

In this problem we will analyze the cache complexity of a couple of algorithms for rotating square matrices. We will assume that all matrices are stored in row-major order.

3(a) [5 Points] Analyze the cache complexity of ITER-MATRIX-ROTATE given in Figure 3.

```

ITER-MATRIX-ROTATE( X, Y, n )
(Input is an  $n \times n$  square matrix  $X[1 : n, 1 : n]$ . This function generates the rotation of  $X$  in  $Y$ .)
1. for  $i \leftarrow 1$  to  $n$  do
2.   for  $j \leftarrow 1$  to  $n$  do
3.      $Y[i, j] \leftarrow X[n - j + 1, i]$ 

```

Figure 3: Iterative matrix rotation.

Solution. Let us assume for simplicity that each cache block can hold B entries of a matrix, and $M \geq 2B$ (i.e., one cache block for writing to Y , and at least one block for reading from X).

Observe that the algorithm accesses X row-by-row and Y column-by-column. Since Y is laid out in row-major order, writing to Y will incur only $\Theta\left(1 + \frac{n^2}{B}\right)$ cache misses in total as long as there is one block dedicated to Y . Since X is also laid out in row-major order but accessed column-by-column, reading the entire matrix will incur $\Theta(n^2)$ cache misses if there are fewer than n cache blocks dedicated to reading X . However, if there are at least n cache blocks for X , then reading entire X will incur only $\mathcal{O}\left(n + \frac{n^2}{B}\right)$ cache misses. Moreover, reading X can incur as few as $\Theta\left(1 + \frac{n^2}{B}\right)$ misses provided at least $2n$ cache blocks are dedicated to the task. Overall, the algorithm will incur $\Theta\left(1 + \frac{n^2}{B} + n^2\right) = \Theta(n^2)$ cache misses provided the cache has fewer than $n + 1$ cache blocks, and $\mathcal{O}\left(1 + \frac{n^2}{B} + n + \frac{n^2}{B}\right) = \mathcal{O}\left(n + \frac{n^2}{B}\right)$ cache misses provided there are between $n + 1$ and $2n$ cache blocks, and $\Theta\left(1 + \frac{n^2}{B}\right)$ misses otherwise.

Hence, the number of cache misses incurred by the algorithm for $n \times n$ input matrices is:

$$Q_{iter}(n) = \begin{cases} \Theta\left(1 + \frac{n^2}{B}\right) & \text{if } M \geq (2n + 1)B, \\ \mathcal{O}\left(n + \frac{n^2}{B}\right) & \text{if } (2n + 1)B > M \geq (n + 1)B, \\ \Theta(n^2) & \text{otherwise.} \end{cases}$$

- 3(b) [**10 Points**] Complete the recursive divide-and-conquer algorithm (REC-MATRIX-ROTATE) for rotating a square matrix given in Figure 4. Analyze its cache complexity assuming a *tall* cache (i.e., $M = \Omega(B^2)$, where M is the cache size and B is the cache block size).

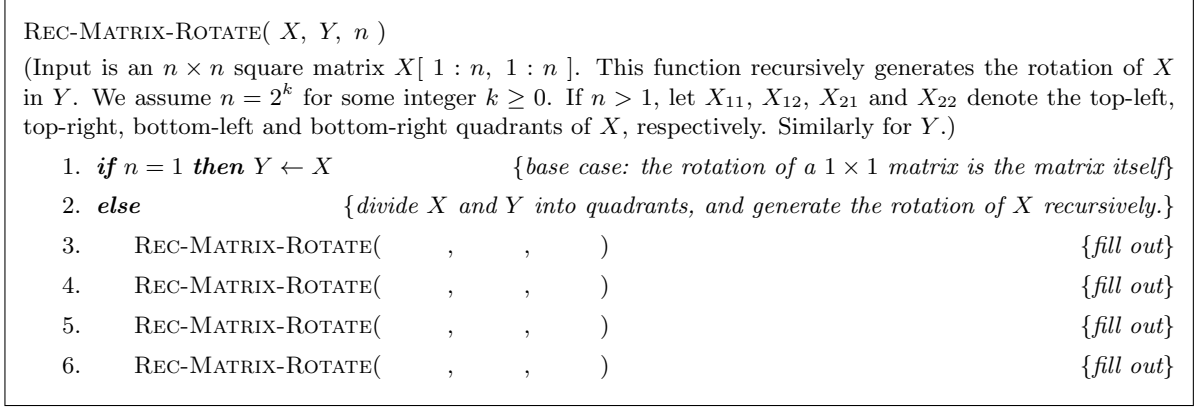


Figure 4: Recursive matrix rotation.

Solution. REC-MATRIX-ROTATE can be completed as follows.

3. REC-MATRIX-ROTATE($X_{11}, Y_{12}, \frac{n}{2}$)
4. REC-MATRIX-ROTATE($X_{12}, Y_{22}, \frac{n}{2}$)
5. REC-MATRIX-ROTATE($X_{21}, Y_{11}, \frac{n}{2}$)
6. REC-MATRIX-ROTATE($X_{22}, Y_{21}, \frac{n}{2}$)

Let $Q(n)$ be the number of cache misses incurred by the algorithm for $n \times n$ input matrices. Then for some suitable constant γ ,

$$Q(n) = \begin{cases} \mathcal{O}\left(n + \frac{n^2}{B}\right) & \text{if } n^2 \leq \gamma M, \\ 4Q\left(\frac{n}{2}\right) & \text{otherwise.} \end{cases}$$

Let $n^2 > \gamma M$ and let k be the smallest positive integer such that $\left(\frac{n}{2^k}\right)^2 \leq \gamma M$. Then expanding the recurrence for $Q(n)$, we obtain:

$$Q(n) = 4^k Q\left(\frac{n}{2^k}\right) = \mathcal{O}\left(\left(\frac{n^2}{M}\right) \left(\sqrt{M} + \frac{M}{B}\right)\right) = \mathcal{O}\left(\frac{n^2}{\sqrt{M}} + \frac{n^2}{B}\right) = \mathcal{O}\left(\frac{n^2}{B}\right) \text{ for } M = \Omega(B^2).$$

However, if the matrices are small enough to fit into the cache then $Q(n) = \mathcal{O}\left(1 + \frac{n^2}{B}\right)$.

Combining the two cases above, we obtain: $Q(n) = \mathcal{O}\left(\frac{n^2}{B} + 1 + \frac{n^2}{B}\right) = \mathcal{O}\left(1 + \frac{n^2}{B}\right)$.

3(c) [**5 Points**] Is the cache complexity result of part 3(b) optimal? Why or why not?

Solution. Observe that any algorithm that tries to rotate matrix X must read X at least once. Since X is an $n \times n$ matrix stored in n^2 contiguous memory locations, one must incur $\Theta\left(1 + \frac{n^2}{B}\right)$ cache misses for reading X once. So, no algorithm for rotating X can incur fewer than $\Theta\left(1 + \frac{n^2}{B}\right)$ cache misses. Since the algorithm in part 3(b) incurs only $\mathcal{O}\left(1 + \frac{n^2}{B}\right)$ cache misses it shows optimal cache performance.

Use this page if you need additional space for your answers.

APPENDIX I: SOME ELEMENTARY PROBABILITY RESULTS

Given an event A , $\Pr[A]$ denotes the probability of occurrence of A . By \overline{A} we denote the opposite or complement of event A . Then $\Pr[\overline{A}]$ denotes the probability of event A not occurring. Clearly,

$$0 \leq \Pr[A], \Pr[\overline{A}] \leq 1 \quad \text{and} \quad \Pr[\overline{A}] = 1 - \Pr[A].$$

Given two events A and B ,

- $A \cap B$ is the event of both A and B occurring, and
- $A \cup B$ is the event of at least one of A and B occurring.

Then the corresponding complements are as follows:

$$\overline{A \cap B} = \overline{A} \cup \overline{B} \quad \text{and} \quad \overline{A \cup B} = \overline{A} \cap \overline{B}.$$

If A and B are mutually exclusive (i.e., both cannot occur simultaneously³), then $\Pr[A \cap B] = 0$. You might find the following relationship useful:

$$\Pr[A \cup B] = \Pr[A] + \Pr[B] - \Pr[A \cap B].$$

Observe that if A and B are mutually exclusive, the relationship given above reduces to:

$$\Pr[A \cup B] = \Pr[A] + \Pr[B].$$

³e.g., if A is the event $(x < 5)$ and B is the event $(x > 5)$ then both A and B cannot be true (i.e., cannot occur) at the same time

APPENDIX II: USEFUL TAIL BOUNDS

Markov's Inequality. Let X be a random variable that assumes only nonnegative values. Then for all $\delta > 0$, $Pr[X \geq \delta] \leq \frac{E[X]}{\delta}$.

Chebyshev's Inequality. Let X be a random variable with a finite mean $E[X]$ and a finite variance $Var[X]$. Then for any $\delta > 0$, $Pr[|X - E[X]| \geq \delta] \leq \frac{Var[X]}{\delta^2}$.

Chernoff Bounds. Let X_1, \dots, X_n be independent Poisson trials, that is, each X_i is a 0-1 random variable with $Pr[X_i = 1] = p_i$ for some p_i . Let $X = \sum_{i=1}^n X_i$ and $\mu = E[X]$. Following bounds hold:

Lower Tail:

- for $0 < \delta < 1$, $Pr[X \leq (1 - \delta)\mu] \leq \left(\frac{e^{-\delta}}{(1-\delta)^{(1-\delta)}}\right)^\mu$
- for $0 < \delta < 1$, $Pr[X \leq (1 - \delta)\mu] \leq e^{-\frac{\mu\delta^2}{2}}$
- for $0 < \gamma < \mu$, $Pr[X \leq \mu - \gamma] \leq e^{-\frac{\gamma^2}{2\mu}}$

Upper Tail:

- for any $\delta > 0$, $Pr[X \geq (1 + \delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$
- for $0 < \delta < 1$, $Pr[X \geq (1 + \delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$
- for $0 < \gamma < \mu$, $Pr[X \geq \mu + \gamma] \leq e^{-\frac{\gamma^2}{3\mu}}$

APPENDIX III: THE MASTER THEOREM

Let $a \geq 1$ and $b > 1$ be constants, let $f(n)$ be a function, and let $T(n)$ be defined on the nonnegative integers by the recurrence

$$T(n) = \begin{cases} \Theta(1), & \text{if } n \leq 1, \\ aT\left(\frac{n}{b}\right) + f(n), & \text{otherwise,} \end{cases}$$

where, $\frac{n}{b}$ is interpreted to mean either $\lfloor \frac{n}{b} \rfloor$ or $\lceil \frac{n}{b} \rceil$. Then $T(n)$ has the following bounds:

Case 1: If $f(n) = \mathcal{O}(n^{\log_b a - \epsilon})$ for some constant $\epsilon > 0$, then $T(n) = \Theta(n^{\log_b a})$.

Case 2: If $f(n) = \Theta(n^{\log_b a} \log^k n)$ for some constant $k \geq 0$, then $T(n) = \Theta(n^{\log_b a} \log^{k+1} n)$.

Case 3: If $f(n) = \Omega(n^{\log_b a + \epsilon})$ for some constant $\epsilon > 0$, and $af\left(\frac{n}{b}\right) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n , then $T(n) = \Theta(f(n))$.