(Lecture 11) Randomized Algorithms and High Probability Bounds

Slide 1

CSE 548: Analysis of Algorithms

Lecture 11
(Randomized Algorithms
& High Probability Bounds)

Rezaul A. Chowdhury

Department of Computer Science
SUNY Stony Brook
Fall 2019

Slide 2

Markov's Inequality

Theorem 1: Let X be a random variable that assumes only nonnegative values. Then for all $\delta > 0$,

$$\Pr[X \geq \delta] \leq \frac{E[X]}{\delta}.$$

Proof: For $\delta > 0$, let

$$I = \begin{cases} 1 & \text{if } X \geq \delta; \\ 0 & \text{otherwise.} \end{cases}$$

Since $X \geq 0$, $I \leq \frac{X}{\delta}$.

We also have, $E[I] = \Pr[I = 1] = \Pr[X \geq \delta]$.

Then $\Pr[X \geq \delta] = E[I] \leq E\left[\frac{X}{\delta}\right] \leq \frac{E[X]}{\delta}$.

Slide 3

Example: Coin Flipping

Let us bound the probability of obtaining more than $\frac{3n}{4}$ heads in a sequence of n fair coin flips.

Let

$$X_i = \begin{cases} 1 & \text{if the } i\text{th coin flip is heads;} \\ 0 & \text{otherwise.} \end{cases}$$

Then the number of heads in n flips, $X = \sum_{i=1}^{n} X_i$.

We know, $E[X_i] = \Pr[X_i = 1] = \frac{1}{2}$.

Hence, $E[X] = \sum_{i=1}^{n} E[X_i] = \frac{n}{2}$.

Then applying Markov's inequality,

$$\Pr\left[X \geq \frac{3n}{4}\right] \leq \frac{E[X]}{3n/4} = \frac{n/2}{3n/4} = \frac{2}{3}.$$

Slide 4

Chebyshev's Inequality

Theorem 2: For any $\delta > 0$,

$$\Pr[|X - E[X]| \geq \delta] \leq \frac{Var[X]}{\delta^2}.$$

Proof: Observe that $\Pr[|X - E[X]| \geq \delta] = \Pr[(X - E[X])^2 \geq \delta^2]$.

Since $(X - E[X])^2$ is a nonnegative random variable, we can use Markov's inequality,

$$\Pr[(X - E[X])^2 \geq \delta^2] \leq \frac{E[(X-E[X])^2]}{\delta^2} = \frac{Var[X]}{\delta^2}.$$

Example: n Fair Coin Flips

$$X_i = \begin{cases} 1 & \text{if the } i\text{th coin flip is heads;} \\ 0 & \text{otherwise.} \end{cases}$$

Then the number of heads in n flips, $X = \sum_{i=1}^{n} X_i$.

We know, $E[X_i] = \Pr[X_i = 1] = \frac{1}{2}$ and $E[(X_i)^2] = E[X_i] = \frac{1}{2}$.

Then $Var[X_i] = E[(X_i)^2] - (E[X_i])^2 = \frac{1}{2} - \frac{1}{4} = \frac{1}{4}$.

Hence, $E[X] = \sum_{i=1}^{n} E[X_i] = \frac{n}{2}$ and $Var[X] = \sum_{i=1}^{n} Var[X_i] = \frac{n}{4}$.

Then applying Chebyshev's inequality,

$$\Pr\left[X \geq \frac{3n}{4}\right] \leq \Pr\left[|X - E[X]| \geq \frac{n}{4}\right] \leq \frac{Var[X]}{(n/4)^2} = \frac{n/4}{(n/4)^2} = \frac{4}{n}.$$

Coin Flipping and Randomized Algorithms

Suppose we have an algorithm that is

– correct (heads) only with probability $p \in (0,1)$, and

– incorrect (tails) with probability $1 - p$.

Question: How many times should we run the algorithm to be reasonably confident that it returns at least one correct solution?

– Las Vegas Algorithm: You keep running the algorithm until you get a correct solution. What is the bound on running time?

– Monte Carlo Algorithm: You stop after a certain number of iterations no matter whether you found a correct solution or not. What is the probability that your solution is correct (or you found a solution)?

Coin Flipping and Randomized Algorithms

Suppose we have an algorithm that is

– correct (heads) only with probability $p \in (0,1)$, and

– incorrect (tails) with probability $1 - p$.

Suppose we run the algorithm k times.

Then probability that no run produces a correct solution is $(1 - p)^k$.

∴ probability of getting at least one correct solution is $1 - (1 - p)^k$.

Set $k = \ln_{\frac{1}{1-p}}\left(\frac{n^\alpha}{c}\right)$, where $\alpha \geq 1$ and $c > 0$ is a constant.

Then the probability that at least one run produces a correct solution is $1 - (1 - p)^k = 1 - \frac{c}{n^\alpha}$.

An event Π is said to occur $\underbrace{\text{with high probability}}_{\text{w.h.p.}}$ if $\Pr[\Pi] \geq 1 - \frac{c}{n^\alpha}$.

Example: A Coloring Problem

Let S be a set of n items.

For $1 \leq l \leq k$, let $S_l \subseteq S$ such that for every pair of $i, j \in [1, k]$ with $i \neq j$, $S_i \neq S_j$ but not necessarily $S_i \cap S_j = \emptyset$.

For each $l \in [1, k]$, let $|S_l| = r$, where $k \leq 2^{r-2}$.

Problem: Color each item of S with one of two colors, red and blue, such that each S_l contains at least one red and one blue item.

Algorithm: Take each item of S and color it either red or blue independently at random (with probability $\frac{1}{2}$).

Clearly, the algorithm does not always lead to a valid coloring (i.e., satisfy the constraints given in our problem statement). What is the probability that it produces a valid coloring?

Example: A Coloring Problem

For $1 \leq l \leq k$, let R_l and B_l be the events that all items of S_l are colored red and blue, respectively.

Then $\Pr[R_l] = \Pr[B_l] = \left(\frac{1}{2}\right)^r = 2^{-r}$ for every $l \in [1, k]$.

$\therefore \Pr\left[\bigcup_{l=1}^{k} R_l\right] = \Pr\left[\bigcup_{l=1}^{k} B_l\right] \leq k 2^{-r} \leq 2^{r-2} 2^{-r} = \frac{1}{4}$.

Thus $\Pr\left[\bigcup_{l=1}^{k}(R_l \cup B_l)\right] \leq 2 \times \frac{1}{4} = \frac{1}{2}$.

$\therefore \Pr\left[\bigcap_{l=1}^{k}(\overline{R_l} \cap \overline{B_l})\right] = 1 - \Pr\left[\bigcup_{l=1}^{k}(R_l \cup B_l)\right] \geq 1 - \frac{1}{2} = \frac{1}{2}$.

Hence, the algorithm is correct with probability at least $\frac{1}{2}$.

To check if the algorithm has produced a correct result we simply check the items in each S_l to verify that neither R_l nor B_l holds.

Hence, we can use this simple algorithm to design a Las Vegas algorithm for solving the coloring problem!

Example: The Min-Cut Problem

Let $G = (V, E)$ be a connected, undirected multigraph with $|V| = n$.

A *cut* in G is a set $C \subseteq E$, such that $G' = (V, E \setminus C)$ is not connected.

A *min-cut* is a cut of minimum cardinality.

The multigraph on the right has a min-cut of size 2: $\{(a, e), (b, c)\}$ and $\{(c, d), (d, e)\}$.

Most deterministic algorithms for finding min-cuts are based on network flows and hence are quite complicated.

Instead in this lecture we will look at a very simple probabilistic algorithm that finds min-cuts with some probability $p > 0$.

Example: The Min-Cut Problem

We apply the following contraction step $n - 2$ times on $G = (V, E)$:

Select an edge (say, (u, v)) from E uniformly at random.

Merge u and v into a single super vertex.

Remove all edges between u and v from E.

If as a result of the contraction there are more than one edges between some pairs of super vertices retain them all.

Let the initial graph be $G_0 = (V_0, E_0)$, where $V_0 = V$ and $E_0 = E$.

Let $G_i = (V_i, E_i)$ be the multigraph after step $i \in [1, n-2]$.

Then clearly, $|V_i| = n - i$ and thus $|V_{n-2}| = 2$.

We return E_{n-2} as our solution.

Example: The Min-Cut Problem

Let us fix our attention on a particular min-cut C of G.

What is the probability that $E_{n-2} = C$?

Suppose $|C| = k$.

Then each vertex of $G_0 = G$ must have degree at least k as otherwise G_0 can be disconnected by removing fewer than k edges.

Hence, $|E| = |E_0| \geq k|V_0|/2 = kn/2$.

Let Π_i be the event of not picking an edge of C for contraction in step $i \in [1, n-2]$.

Then clearly, $\Pr[\Pi_1] = 1 - \frac{k}{|E_0|} \geq 1 - \frac{k}{kn/2} = 1 - \frac{2}{n}$

Also $\Pr[\Pi_2 | \Pi_1] = 1 - \frac{k}{|E_1|} \geq 1 - \frac{k}{k(n-1)/2} = 1 - \frac{2}{n-1}$

In general, $\Pr\left[\Pi_i | \bigcap_{j=1}^{i-1} \Pi_j\right] = 1 - \frac{k}{|E_{i-1}|} \geq 1 - \frac{k}{k(n-i+1)/2} = 1 - \frac{2}{n-i+1}$

(Lecture 11) Randomized Algorithms and High Probability Bounds

Example: The Min-Cut Problem

The probability that no edge of C was ever picked by the algorithm is:

$$\Pr[\cap_{i=1}^{n-2} \Pi_i] \geq \prod_{i=1}^{n-2} \left(1 - \frac{2}{n-i+1}\right) = \frac{2}{n(n-1)} > \frac{2}{n^2}.$$

So $\Pr[E_{n-2} = C] > \frac{2}{n^2}$, and $\Pr[E_{n-2} \neq C] < 1 - \frac{2}{n^2}$.

Suppose we run the algorithm $n^2/2$ times, and return the smallest cut, say C', obtained from those $n^2/2$ attempts.

Then $\Pr[C' \neq C] < \left(1 - \frac{2}{n^2}\right)^{n^2/2} < \frac{1}{e} \Rightarrow \Pr[C' = C] > 1 - \frac{1}{e}$.

Hence, the algorithm will return a min-cut with probability $> 1 - \frac{1}{e}$.

But we do not know how to detect if the cut returned by the algorithm is, indeed, a min-cut.

Still we can design a Monte-Carlo algorithm based on this simple idea to produce a min-cut with high probability!

When Only One Success is Not Enough

In both examples we have looked at so far, we were happy with only one success. The analysis was easy.

But sometimes we need the algorithm to be successful for at least or at most a certain number of times (we will see a very familiar such example shortly).

The number of successful runs required often depends on the size of the input.

How do we analyze those algorithms?

Preparing for Chernoff Bounds

Lemma 1: Let $X_1, \ldots, X_n$ be independent Poisson trials, that is, each X_i is a 0-1 random variable with $\Pr[X_i = 1] = p_i$ for some p_i. Let $X = \sum_{i=1}^{n} X_i$ and $\mu = E[X]$. Then for any $t > 0$,

$$E[e^{tX}] \leq e^{(e^t - 1)\mu}.$$

Proof: $E[e^{tX_i}] = p_i e^{t \times 1} + (1 - p_i)e^{t \times 0} = p_i e^t + (1 - p_i)$
$$= 1 + p_i(e^t - 1)$$

But for any y, $1 + y \leq e^y$. Hence, $E[e^{tX_i}] \leq e^{p_i(e^t - 1)}$.

Now, $E[e^{tX}] = E\left[e^{t \sum_{i=1}^{n} X_i}\right] = E\left[\prod_{i=1}^{n} e^{tX_i}\right] = \prod_{i=1}^{n} E[e^{tX_i}]$

$$\leq \prod_{i=1}^{n} e^{p_i(e^t - 1)} = e^{(e^t - 1)\sum_{i=1}^{n} p_i}$$

But, $\mu = E[X] = E\left[\sum_{i=1}^{n} X_i\right] = \sum_{i=1}^{n} E[X_i] = \sum_{i=1}^{n} p_i$.

Hence, $E[e^{tX}] \leq e^{(e^t - 1)\mu}$

Chernoff Bound 1

Theorem 3: Let $X_1, \ldots, X_n$ be independent Poisson trials, that is, each X_i is a 0-1 random variable with $\Pr[X_i = 1] = p_i$ for some p_i. Let $X = \sum_{i=1}^{n} X_i$ and $\mu = E[X]$. Then for any $\delta > 0$,

$$\Pr[X \geq (1 + \delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu.$$

Proof: Applying Markov's inequality for any $t > 0$,

$$\Pr[X \geq (1 + \delta)\mu] = \Pr[e^{tX} \geq e^{t(1+\delta)\mu}] \leq \frac{E[e^{tX}]}{e^{t(1+\delta)\mu}}$$

$$\leq \frac{e^{(e^t - 1)\mu}}{e^{t(1+\delta)\mu}} \quad [\text{Lemma 1}]$$

Setting $t = \ln(1 + \delta) > 0$, i.e., $e^t = 1 + \delta$, we get,

$$\Pr[X \geq (1 + \delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu.$$

Chernoff Bound 2

17

Theorem 4: For $0 < \delta < 1$, $\Pr[X \geq (1+\delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$.

Proof: From Theorem 3, for $\delta > 0$, $\Pr[X \geq (1+\delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$.

We will show that for $0 < \delta < 1$, $\frac{e^\delta}{(1+\delta)^{(1+\delta)}} \leq e^{-\frac{\delta^2}{3}}$

$$\Rightarrow \delta - (1+\delta)\ln(1+\delta) \leq -\frac{\delta^2}{3}$$

That is, $f(\delta) = \delta - (1+\delta)\ln(1+\delta) + \frac{\delta^2}{3} \leq 0$

We have, $f'(\delta) = -\ln(1+\delta) + \frac{2}{3}\delta$, and $f''(\delta) = -\frac{1}{1+\delta} + \frac{2}{3}$

Observe that $f''(\delta) < 0$ for $0 \leq \delta \leq \frac{1}{2}$, and $f''(\delta) > 0$ for $\delta > \frac{1}{2}$.

Chernoff Bound 2

18

Theorem 4: For $0 < \delta < 1$, $\Pr[X \geq (1+\delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$.

Proof: From Theorem $+\delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$.

We will show that fo $e^{-\frac{\delta^2}{3}}$

$\leq -\frac{\delta^2}{3}$

That is, $f(\delta) = \delta -$ 0

We have, $f'(\delta) = -\ln(1+\delta) + \frac{2}{3}\delta$, and $f''(\delta) = -\frac{1}{1+\delta} + \frac{2}{3}$

Observe that $f''(\delta) < 0$ for $0 \leq \delta \leq \frac{1}{2}$, and $f''(\delta) > 0$ for $\delta > \frac{1}{2}$.

Chernoff Bound 2

19

Theorem 4: For $0 < \delta < 1$, $\Pr[X \geq (1+\delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$.

Proof: From Theorem 3, for $\delta > 0$, $\Pr[X \geq (1+\delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$.

We will show that for $0 < \delta < 1$, $\frac{e^\delta}{(1+\delta)^{(1+\delta)}} \leq e^{-\frac{\delta^2}{3}}$

$$\Rightarrow \delta - (1+\delta)\ln(1+\delta) \leq -\frac{\delta^2}{3}$$

That is, $f(\delta) = \delta - (1+\delta)\ln(1+\delta) + \frac{\delta^2}{3} \leq 0$

We have, $f'(\delta) = -\ln(1+\delta) + \frac{2}{3}\delta$, and $f''(\delta) = -\frac{1}{1+\delta} + \frac{2}{3}$

Observe that $f''(\delta) < 0$ for $0 \leq \delta \leq \frac{1}{2}$, and $f''(\delta) > 0$ for $\delta > \frac{1}{2}$.

Hence, $f'(\delta)$ first decreases and then increases over $[0,1]$.

Since $f'(0) = 0$ and $f'(1) < 0$, we have $f'(\delta) \leq 0$ over $[0,1]$.

Chernoff Bound 2

20

Theorem 4: For $0 < \delta < 1$, $\Pr[X \geq (1+\delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$.

Proof: From Theorem $+\delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$.

We will show that fo $e^{-\frac{\delta^2}{3}}$

$\leq -\frac{\delta^2}{3}$

That is, $f(\delta) = \delta -$ 0

We have, $f'(\delta) = -\ln(1+\delta) + \frac{2}{3}\delta$, and $f''(\delta) = -\frac{1}{1+\delta} + \frac{2}{3}$

Observe that $f''(\delta) < 0$ for $0 \leq \delta \leq \frac{1}{2}$, and $f''(\delta) > 0$ for $\delta > \frac{1}{2}$.

Hence, $f'(\delta)$ first decreases and then increases over $[0,1]$.

Since $f'(0) = 0$ and $f'(1) < 0$, we have $f'(\delta) \leq 0$ over $[0,1]$.

Chernoff Bound 2

21

Theorem 4: For $0 < \delta < 1$, $\Pr[X \geq (1+\delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$.

Proof: From Theorem 3, for $\delta > 0$, $\Pr[X \geq (1+\delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$.

We will show that for $0 < \delta < 1$, $\frac{e^\delta}{(1+\delta)^{(1+\delta)}} \leq e^{-\frac{\delta^2}{3}}$

$$\Rightarrow \delta - (1+\delta)\ln(1+\delta) \leq -\frac{\delta^2}{3}$$

That is, $f(\delta) = \delta - (1+\delta)\ln(1+\delta) + \frac{\delta^2}{3} \leq 0$

We have, $f'(\delta) = -\ln(1+\delta) + \frac{2}{3}\delta$, and $f''(\delta) = -\frac{1}{1+\delta} + \frac{2}{3}$

Observe that $f''(\delta) < 0$ for $0 \leq \delta \leq \frac{1}{2}$, and $f''(\delta) > 0$ for $\delta > \frac{1}{2}$.

Hence, $f'(\delta)$ first decreases and then increases over $[0,1]$.

Since $f'(0) = 0$ and $f'(1) < 0$, we have $f'(\delta) \leq 0$ over $[0,1]$.

Since $f(0) = 0$, it follows that $f(\delta) \leq 0$ in that interval.

Chernoff Bound 2

22

Theorem 4: For $0 < $ [ ] $\leq e^{-\frac{\mu\delta^2}{3}}$.

Proof: From Theorem [ ] $+\delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$.

We will show that fo [ ] $e^{-\frac{\delta^2}{3}}$

[ ] $\leq -\frac{\delta^2}{3}$

That is, $f(\delta) = \delta - (1+\delta)\ln(1+\delta) + \frac{o}{3} \leq 0$

We have, $f'(\delta) = -\ln(1+\delta) + \frac{2}{3}\delta$, and $f''(\delta) = -\frac{1}{1+\delta} + \frac{2}{3}$

Observe that $f''(\delta) < 0$ for $0 \leq \delta \leq \frac{1}{2}$, and $f''(\delta) > 0$ for $\delta > \frac{1}{2}$.

Hence, $f'(\delta)$ first decreases and then increases over $[0,1]$.

Since $f'(0) = 0$ and $f'(1) < 0$, we have $f'(\delta) \leq 0$ over $[0,1]$.

Since $f(0) = 0$, it follows that $f(\delta) \leq 0$ in that interval.

Chernoff Bound 2

23

Theorem 4: For $0 < \delta < 1$, $\Pr[X \geq (1+\delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$.

Proof: From Theorem 3, for $\delta > 0$, $\Pr[X \geq (1+\delta)\mu] \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}}\right)^\mu$.

We will show that for $0 < \delta < 1$, $\frac{e^\delta}{(1+\delta)^{(1+\delta)}} \leq e^{-\frac{\delta^2}{3}}$

$$\Rightarrow \delta - (1+\delta)\ln(1+\delta) \leq -\frac{\delta^2}{3}$$

That is, $f(\delta) = \delta - (1+\delta)\ln(1+\delta) + \frac{\delta^2}{3} \leq 0$

We have, $f'(\delta) = -\ln(1+\delta) + \frac{2}{3}\delta$, and $f''(\delta) = -\frac{1}{1+\delta} + \frac{2}{3}$

Observe that $f''(\delta) < 0$ for $0 \leq \delta \leq \frac{1}{2}$, and $f''(\delta) > 0$ for $\delta > \frac{1}{2}$.

Hence, $f'(\delta)$ first decreases and then increases over $[0,1]$.

Since $f'(0) = 0$ and $f'(1) < 0$, we have $f'(\delta) \leq 0$ over $[0,1]$.

Since $f(0) = 0$, it follows that $f(\delta) \leq 0$ in that interval.

Chernoff Bound 3

24

Corollary 1: For $0 < \gamma < \mu$, $\Pr[X \geq \mu + \gamma] \leq e^{-\frac{\gamma^2}{3\mu}}$.

Proof: From Theorem 2, for $0 < \delta < 1$, $\Pr[X \geq (1+\delta)\mu] < e^{-\frac{\mu\delta^2}{3}}$.

Setting $\gamma = \mu\delta$, we get, $\Pr[X \geq \mu + \gamma] \leq e^{-\frac{\gamma^2}{3\mu}}$ for $0 < \gamma < \mu$.

Example: n Fair Coin Flips

25

$$X_i = \begin{cases} 1 & \text{if the } i\text{th coin flip is heads;} \\ 0 & \text{otherwise.} \end{cases}$$

Then the number of heads in n flips, $X = \sum_{i=1}^{n} X_i$.

We know, $E[X_i] = \Pr[X_i = 1] = \frac{1}{2}$.

Hence, $\mu = E[X] = \sum_{i=1}^{n} E[X_i] = \frac{n}{2}$.

Now putting $\delta = \frac{1}{2}$ in Chernoff bound 2, we have,

$$\Pr\left[X \geq \frac{3n}{4}\right] \leq e^{-\frac{n}{24}} = \frac{1}{e^{\frac{n}{24}}}.$$

Chernoff Bounds 4, 5 and 6

26

Theorem 5: For $0 < \delta < 1$, $\Pr[X \leq (1 - \delta)\mu] \leq \left(\frac{e^{-\delta}}{(1-\delta)^{(1-\delta)}}\right)^{\mu}$.

Theorem 6: For $0 < \delta < 1$, $\Pr[X \leq (1 - \delta)\mu] \leq e^{-\frac{\mu\delta^2}{2}}$.

Corollary 2: For $0 < \gamma < \mu$, $\Pr[X \leq \mu - \gamma] \leq e^{-\frac{\gamma^2}{2\mu}}$.

Chernoff Bounds

27

Lower Tail	Upper Tail
$0 < \delta < 1$: $\Pr[X \leq (1 - \delta)\mu] \leq \left(\frac{e^{-\delta}}{(1-\delta)^{(1-\delta)}}\right)^{\mu}$	$\delta > 0$: $\Pr[X \geq (1 + \delta)\mu] \leq \left(\frac{e^{\delta}}{(1+\delta)^{(1+\delta)}}\right)^{\mu}$
$0 < \delta < 1$: $\Pr[X \leq (1 - \delta)\mu] \leq e^{-\frac{\mu\delta^2}{2}}$	$0 < \delta < 1$: $\Pr[X \geq (1 + \delta)\mu] \leq e^{-\frac{\mu\delta^2}{3}}$
$0 < \gamma < \mu$: $\Pr[X \leq \mu - \gamma] \leq e^{-\frac{\gamma^2}{2\mu}}$	$0 < \gamma < \mu$: $\Pr[X \geq \mu + \gamma] \leq e^{-\frac{\gamma^2}{3\mu}}$

$\Pr[Y \approx x]$

$e^{-\Theta(\varepsilon^2\mu)}$

$e^{-\Theta(1)}$

$n^{-\Theta(1)}$

$\sqrt{\mu}$

$\varepsilon\mu$

$\sqrt{\mu \ln n}$

$\mu = E[Y]$

Source
greedyalgs.info

28

Randomized Quicksort
# (RANDQS)

Randomized Quicksort (RANDQS)

29

Input: A set of numbers S. (i.e., all numbers are distinct)

Output: The numbers of S sorted in increasing order.

Steps:

1. **Pivot Selection:** Select a number $x \in S$ uniformly at random.

2. **Partition:** Compare each number of S with x, and determine sets $S_l = \{y \in S \mid y < x\}$ and $S_r = \{y \in S \mid y > x\}$.

3. **Recursion:** Recursively sort S_l and S_r.

4. **Output:** Output the sorted version of S_l, followed by x, followed by the sorted version of S_r.

Randomized Quicksort (RANDQS)

30

Input: A set of numbers S. (i.e., all numbers are distinct)

Output: The numbers of S sorted in increasing order.

Steps:

1. **Pivot Selection:** Select a number $x \in S$ uniformly at random.

2. **Partition:** Compare each number of S with x, and determine sets $S_l = \{y \in S \mid y < x\}$ and $S_r = \{y \in S \mid y > x\}$.

3. **Recursion:** Recursively sort S_l and S_r.

4. **Output:** Output the sorted version of S_l, followed by x, followed by the sorted version of S_r.

Assumption: RANDQS is called only on nonempty S.

Observation: If $|S| = n$, fewer than n recursive calls to RANDQS will be made during the sorting of S. (why?)

Observation: If $|S| = n$, and X is the total number of comparisons made in step 2 (Partition) across all (original and recursive) calls to RANDQS, then RANDQS sorts S in $\mathrm{O}(n + X)$ time.

Expected Running Time of RANDQS

31

Observation: If $|S| = n$, and X is the total number of comparisons made in step 2 (Partition) across all (original and recursive) calls to RANDQS, then RANDQS sorts S in $\mathrm{O}(n + X)$ time.

Then all we need to do is determine $E[X]$.

Let $s_1, s_2, \ldots, s_n$ be the elements of S in sorted order.

Let $S_{ij} = \{s_i, s_{i+1}, \ldots, s_j\}$ for all $1 \le i < j \le n$.

Observe that each pair of elements of S is compared at most once during the entire execution of the algorithm. (why?)

For $1 \le i < j \le n$, let $X_{ij} = \begin{cases} 1 & \text{if } s_i \text{ is compared to } s_j; \\ 0 & \text{otherwise.} \end{cases}$

Then $X = \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} X_{ij}$.

Expected Running Time of RANDQS

32

For $1 \le i < j \le n$, let $X_{ij} = \begin{cases} 1 & \text{if } s_i \text{ is compared to } s_j; \\ 0 & \text{otherwise.} \end{cases}$

Then $X = \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} X_{ij}$

$\Rightarrow E[X] = \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} E[X_{ij}]$

$\qquad = \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} \Pr[X_{ij} = 1]$

Observations:

$X_{ij} = 0$: Once a pivot x with $s_i < x < s_j$ is chosen, s_i and s_j will never be compared at any subsequent time. (why?)

$X_{ij} = 1$: If either s_i or s_j is chosen as a pivot before any other item in S_{ij} then s_i will be compared with s_j. (why?)

(Lecture 11) Randomized Algorithms and High Probability Bounds

Expected Running Time of RANDQS

Since each element of S_{ij} is equally likely to be chosen as a pivot:

$$\Pr[X_{ij} = 1] \leq \frac{1}{j-i+1} + \frac{1}{j-i+1} = \frac{2}{j-i+1}.$$

Hence, $E[X] = \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} \Pr[X_{ij} = 1]$

$$\leq \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} \frac{2}{j-i+1}$$

$$= \sum_{i=1}^{n-1} \sum_{k=1}^{n-i} \frac{2}{k+1}$$

$$< \sum_{i=1}^{n-1} \sum_{k=1}^{n} \frac{2}{k}$$

$$= \sum_{i=1}^{n-1} O(\log n)$$

$$= O(n \log n)$$

Thus expected running time of RANDQS is $O(n \log n)$.

High Probability Bound for RANDQS

Input: A set of numbers S. (i.e., all numbers are distinct)

Output: The numbers of S sorted in increasing order.

Steps:

1. **Pivot Selection:** Select a number $x \in S$ uniformly at random.
2. **Partition:** Compare each number of S with x, and determine sets $S_l = \{y \in S | y < x\}$ and $S_r = \{y \in S | y > x\}$.
3. **Recursion:** Recursively sort S_l and S_r.
4. **Output:** Output the sorted version of S_l, followed by x, followed by the sorted version of S_r.

We will prove that w.h.p. the running time of RANDQS does not exceed its expected running time by more than a constant factor.

In other words, we show that w.h.p. RANDQS runs in $O(n \log n)$ time.

High Probability Bound for RANDQS

Input: A set of numbers S. (i.e., all numbers are distinct)

Output: The numbers of S sorted in increasing order.

Steps:

1. **Pivot Selection:** Select a number $x \in S$ uniformly at random.
2. **Partition:** Compare each number of S with x, and determine sets $S_l = \{y \in S | y < x\}$ and $S_r = \{y \in S | y > x\}$.
3. **Recursion:** Recursively sort S_l and S_r.
4. **Output:** Output the sorted version of S_l, followed by x, followed by the sorted version of S_r.

Let us fix an element z in the original input set of size n.

We will trace the partition containing z for $c \ln n$ levels of recursion, where c is a constant to be determined later.

If a partitioning step divides S such that $\frac{|S|}{4} \leq |S_l|, |S_r| \leq \frac{3|S|}{4}$, we call that partition a *balanced* partition.

High Probability Bound for RANDQS

We will prove that among the $c \ln n$ partitioning steps z undergoes, w.h.p. at least $\frac{c}{4} \ln n$ results in balanced partitions.

If at any point z is in a partition of size k, after a balanced partitioning step it ends up in a partition of size at most $\left(\frac{3}{4}\right) k$.

Since the input size is n, after $\frac{c}{4} \ln n$ balanced partitions, z will end up in a partition of size $\leq \left(\frac{3}{4}\right)^{\frac{c}{4} \ln n} n = \frac{n}{n^{\frac{c}{4} \ln\left(\frac{4}{3}\right)}}$, which is ≤ 1 for $c \geq 14$.

That means if $c \geq 14$, then z will end up in its final sorted position in the output after undergoing $\frac{c}{4} \ln n$ balanced partitions.

High Probability Bound for RANDQS [37]

For $1 \le i \le c \ln n$, let

$$Z_i = \begin{cases} 1 & \text{if the partition at recursion level } i \text{ is balanced;} \\ 0 & \text{otherwise.} \end{cases}$$

But a balanced partition is obtained by choosing a pivot with rank between $\frac{k}{4}$ and $\frac{3k}{4}$, where k is the size of the set being partitioned.

Since each element of the set is chosen as a pivot uniformly at random, a balancing pivot will be chosen with probability $\frac{\frac{3k}{4} - \frac{k}{4}}{k} = \frac{1}{2}$.

Hence, $\Pr[Z_i = 1] = \frac{1}{2}$.

Thus $E[Z_i] = \Pr[Z_i = 1] = \frac{1}{2}$.

High Probability Bound for RANDQS [38]

Total number of balanced partitions, $Z = \sum_{i=1}^{c \ln n} Z_i$.

Then $\mu = E[Z] = \sum_{i=1}^{c \ln n} E[Z_i] = \frac{c \ln n}{2}$.

Now applying Chernoff bound 5 (see Theorem 6) with $\delta = \frac{1}{2}$,

$$\Pr[Z \le (1 - \delta)\mu] \le e^{-\frac{\mu \delta^2}{2}}$$

$$\Rightarrow \Pr\left[Z \le \frac{c}{4} \ln n\right] \le e^{-\frac{\mu \delta^2}{2}} = e^{-\frac{c}{16} \ln n} = n^{-\frac{c}{16}} = \frac{1}{n^{\frac{c}{16}}}.$$

For $c = 32$, we have $\Pr[Z \le 8 \ln n] \le \frac{1}{n^2}$.

This means that the probability that z fails to reach its final sorted position even after $32 \ln n$ levels of recursion is $\le \frac{1}{n^2}$.

High Probability Bound for RANDQS [39]

The probability that at least one of n input elements fails to reach its final sorted position after $32 \ln n$ levels of recursion is $\le n \times \frac{1}{n^2} = \frac{1}{n}$.

$\therefore$ the probability that all n input elements reach their final sorted positions after $32 \ln n$ levels of recursion is $\ge 1 - \frac{1}{n}$.

But observe that the total amount of work done in each level of recursion is $O(n)$.

$\therefore$ total work done in $32 \ln n$ levels of recursion is $O(n \log n)$.

Hence, w.h.p. RANDQS terminates in $O(n \log n)$ time.

Random Skip Lists [40]

(Lecture 11) Randomized Algorithms and High Probability Bounds

Searching in a Sorted Linked List

Traditional Linked List:

SEARCH(x): Takes $\leq n$ time.

2-level Linked List:

SEARCH(x): Takes $\leq 2\sqrt{n}$ time.

Searching in a Sorted Linked List

3-level Linked List:

SEARCH(x): Takes $\leq 3\sqrt[3]{n}$ time.

k-**level Linked List:** SEARCH(x) takes $\leq k\sqrt[k]{n} = kn^{\frac{1}{k}}$ time.

For $k = \log n$**:** SEARCH(x) takes $\leq (\log n) \cdot n^{\frac{1}{\log n}} = 2\log n$ time!

Searching in a Sorted Linked List

$(\log n)$-**level Linked List:** SEARCH takes $\leq (\log n) \cdot n^{\frac{1}{\log n}} = 2\log n$ time!

Observations:

1. Let $n_l =$ #items in level l. Then $n_{l+1} = \left\lceil \frac{n_l}{2} \right\rceil$.

2. Let $m_l = n_l - n_{l+1} =$ #items in level l that have not reached level $l + 1$. Then $m_l = \left\lfloor \frac{n}{2^l} \right\rfloor$.

Searching in a Sorted Linked List

$(\log n)$-**level Linked List:** SEARCH takes $\leq (\log n) \cdot n^{\frac{1}{\log n}} = 2\log n$ time!

How do we maintain this regular structure under insertion and deletion of items?

Deterministic solution does not seem straightforward.

But randomization can make life really easy!

Random Skip Lists

Construction:

1. Start with all items along with a sentinel $-\infty$ in level 1.
2. Promote each non-sentinel item of level $l > 0$ to level $l + 1$ with probability $\frac{1}{2}$.

   If level $l + 1$ is nonempty promote the sentinel, too.

Random Skip Lists

Let L be a skip list,

L_k be the set of all items in level $k \geq 1$,

$l(x) = \max\{ k \mid x \in L_k \}$, and

$h(L) = \max\{ l(x) \mid x \in L_0 \}$.

Random Skip Lists

Clearly, for each $x \in L$ and $k \geq 1$, $\Pr[\, l(x) = k \,] = \frac{1}{2^k}$.

Then $\Pr[\, l(x) > k \,] = \sum_{i=k+1}^{\infty} \Pr[\, l(x) = i \,] = \sum_{i=k+1}^{\infty} \left(\frac{1}{2^i}\right) = \frac{1}{2^k}$.

$\therefore \Pr[\, h(L) \leq k \,] = \prod_{x \in L} \Pr[\, h(L) \leq k \,] = \left(1 - \frac{1}{2^k}\right)^n$.

$\therefore$ For constant $c > 2$, $\Pr[\, h(L) \leq c \log n \,]$

$$= \left(1 - \frac{1}{2^{c \log n}}\right)^n = \left(1 - \frac{1}{n^c}\right)^n \geq 1 - \frac{n}{n^c} = 1 - \frac{1}{n^c}$$

Hence, w.h.p. height of a skip list is $O(\log n)$.

Random Skip Lists

Let us flip $4c \log n$ fair coins, and let X is the number of heads we get.

Then $\mu = E[X] = (4c \log n) \times \frac{1}{2} = 2c \log n$.

We know for $0 < \delta < 1$, Chernoff bound, $\Pr[X \leq (1 - \delta)\mu] \leq e^{-\frac{\mu \delta^2}{2}}$.

Putting $\delta = \frac{1}{2}$ and $\mu = 2c \log n$, we get, $\Pr[X \leq c \log n] \leq \frac{1}{n^{\frac{c}{4}}}$.

For $c \geq 16$, $\Pr[X > c \log n] \geq 1 - \frac{1}{n}$.

Hence, w.h.p. we will get more than $c \log n$ heads.