

1

CSE 548: Analysis of Algorithms

Lecture 12

(Analyzing Parallel Algorithms)

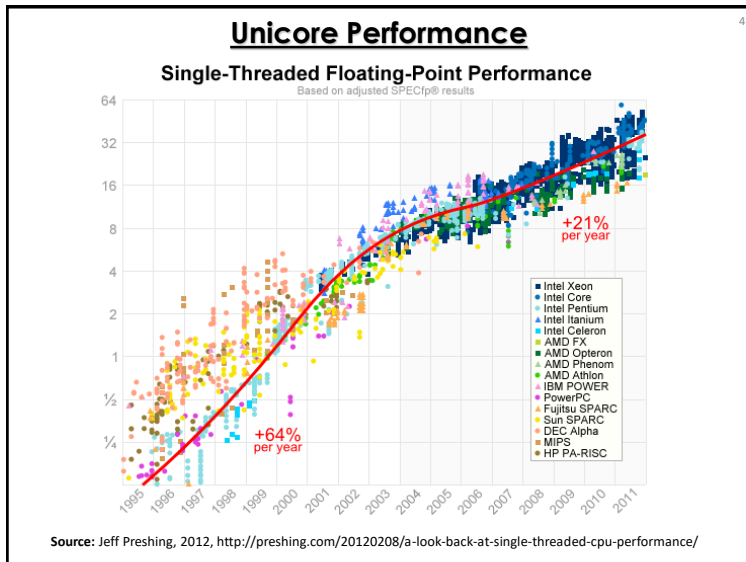
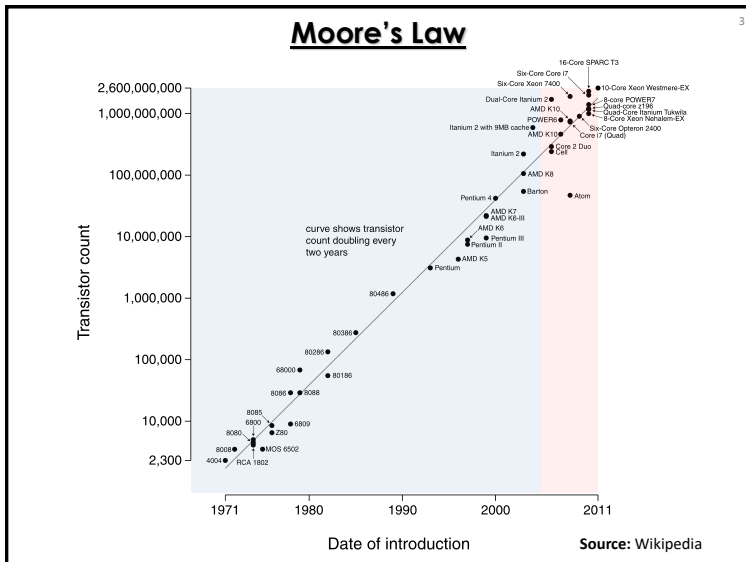
Rezaul A. Chowdhury

Department of Computer Science

SUNY Stony Brook

Fall 2019

Why Parallelism?



Unicore Performance Has Hit a Wall!

Some Reasons

- Lack of additional ILP
(Instruction Level Hidden Parallelism)
- High power density
- Manufacturing issues
- Physical limits
- Memory speed



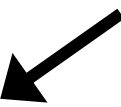
Unicore Performance: No Additional ILP

“Everything that can be invented has been invented.”

— Charles H. Duell
Commissioner, U.S. patent office, 1899

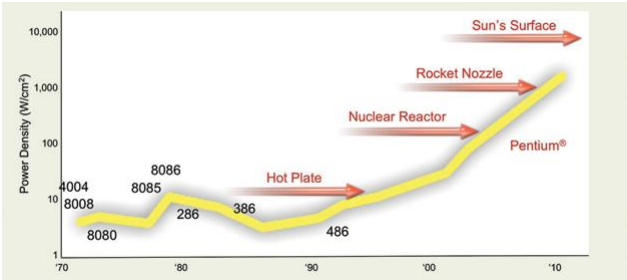
Exhausted all ideas to exploit hidden parallelism?

- Multiple simultaneous instructions
- Instruction Pipelining
- Out-of-order instructions
- Speculative execution
- Branch prediction
- Register renaming, etc.



Unicore Performance: High Power Density

- Dynamic power, $P_d \propto V^2 f C$
 - V = supply voltage
 - f = clock frequency
 - C = capacitance
- But $V \propto f$
- Thus $P_d \propto f^3$



Source: Patrick Gelsinger, Intel Developer Forum, Spring 2004 (Simon Floyd)



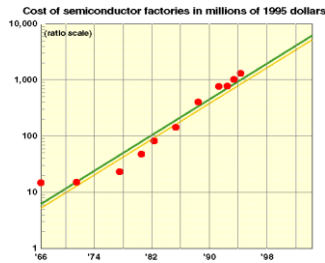
Unicore Performance: Manufacturing Issues

- Frequency, $f \propto 1 / s$
 - s = feature size (transistor dimension)
- Transistors / unit area $\propto 1 / s^2$
- Typically, die size $\propto 1 / s$
- So, what happens if feature size goes down by a factor of x ?
 - Raw computing power goes up by a factor of x^4 !
 - Typically most programs run faster by a factor of x^3 without any change!

Source: Kathy Yelick and Jim Demmel, UC Berkeley

Unicore Performance: Manufacturing Issues

- Manufacturing cost goes up as feature size decreases
 - Cost of a semiconductor fabrication plant doubles every 4 years (Rock's Law)
- CMOS feature size is limited to 5 nm (at least 10 atoms)



Source: Kathy Yelick and Jim Demmel, UC Berkeley

Unicore Performance: Physical Limits

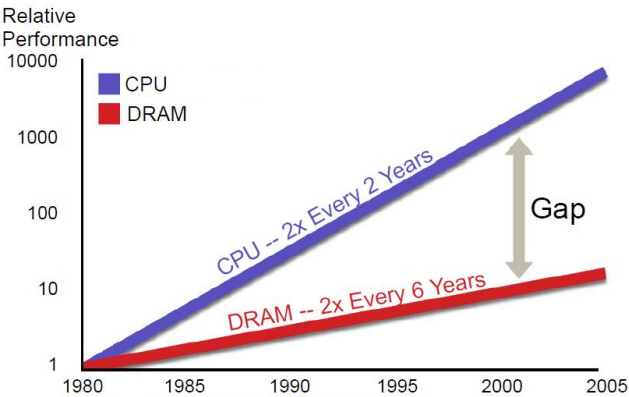
Execute the following loop on a serial machine in 1 second:

```
for ( i = 0; i < 1012; ++i )  
    z[ i ] = x[ i ] + y[ i ];
```

- We will have to access 3×10^{12} data items in one second
- Speed of light is, $c \approx 3 \times 10^8$ m/s
- So each data item must be within $c / 3 \times 10^{12} \approx 0.1$ mm from the CPU on the average
- All data must be put inside a 0.2 mm \times 0.2 mm square
- Each data item (≥ 8 bytes) can occupy only 1 \AA^2 space! (size of a small atom!)

Source: Kathy Yelick and Jim Demmel, UC Berkeley

Unicore Performance: Memory Wall



Source: Sun World Wide Analyst Conference Feb. 25, 2003

Source: Rick Hetherington, Chief Technology Officer, Microelectronics, Sun Microsystems

Unicore Performance Has Hit a Wall!

Some Reasons

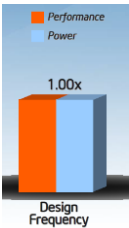
- Lack of additional ILP (Instruction Level Hidden Parallelism)
- High power density
- Manufacturing issues
- Physical limits
- Memory speed

“Oh Sinnerman, where you gonna run to?”

— Sinnerman (recorded by Nina Simone)

Where You Gonna Run To?

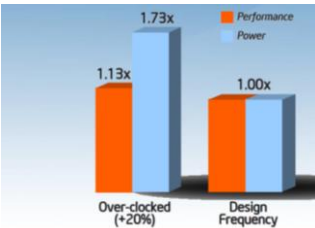
- Changing f by 20% changes performance by 13%
- So what happens if we overclock by 20%?



Source: Andrew A. Chien, Vice President of Research, Intel Corporation

Where You Gonna Run To?

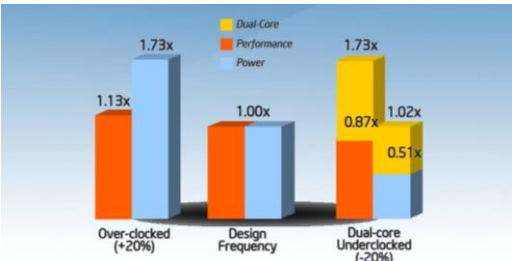
- Changing f by 20% changes performance by 13%
- So what happens if we overclock by 20%?
- And underclock by 20%?



Source: Andrew A. Chien, Vice President of Research, Intel Corporation

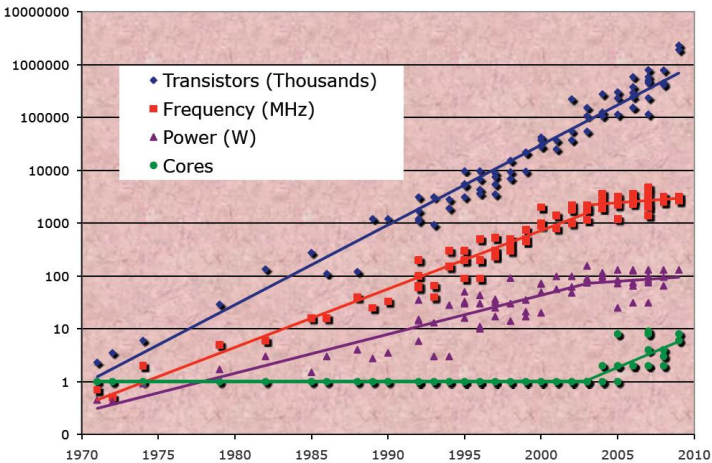
Where You Gonna Run To?

- Changing f by 20% changes performance by 13%
- So what happens if we overclock by 20%?
- And underclock by 20%?

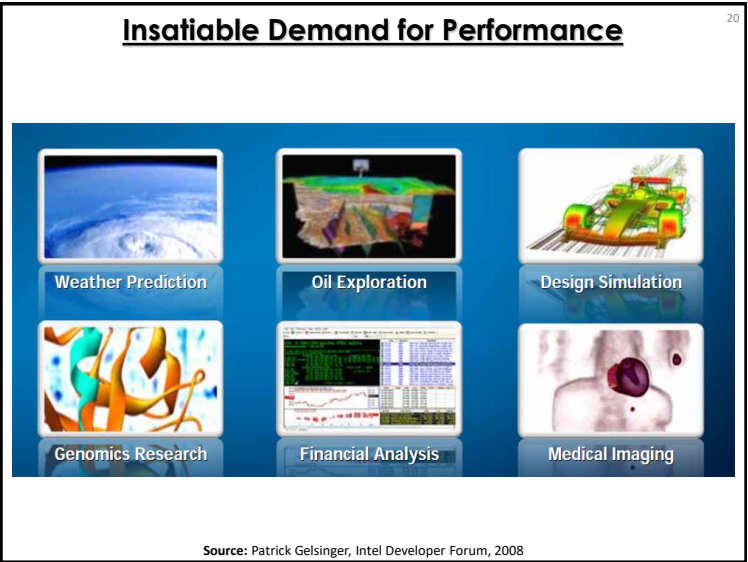
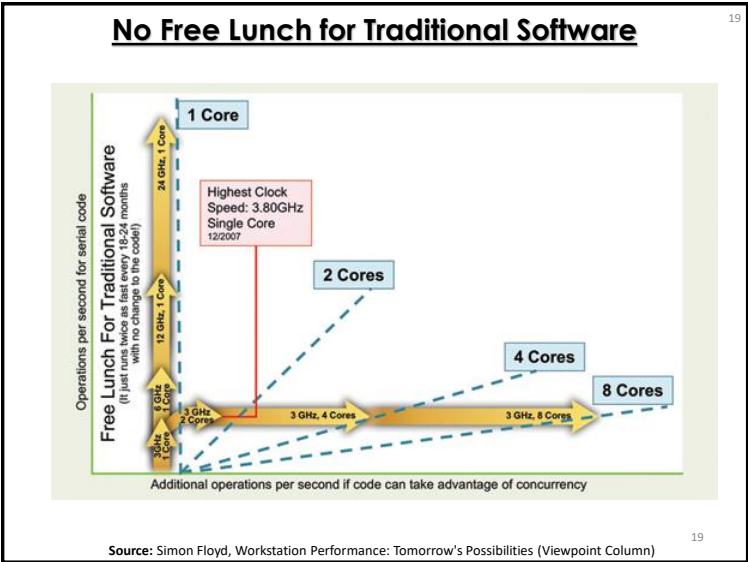
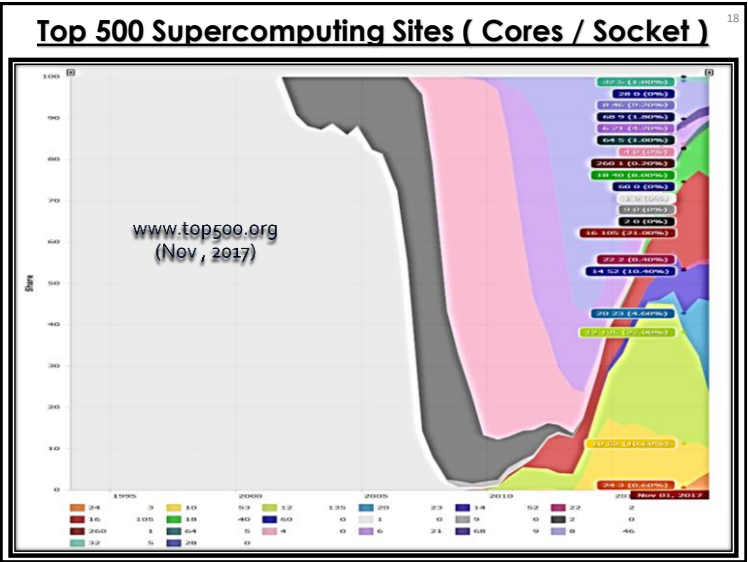
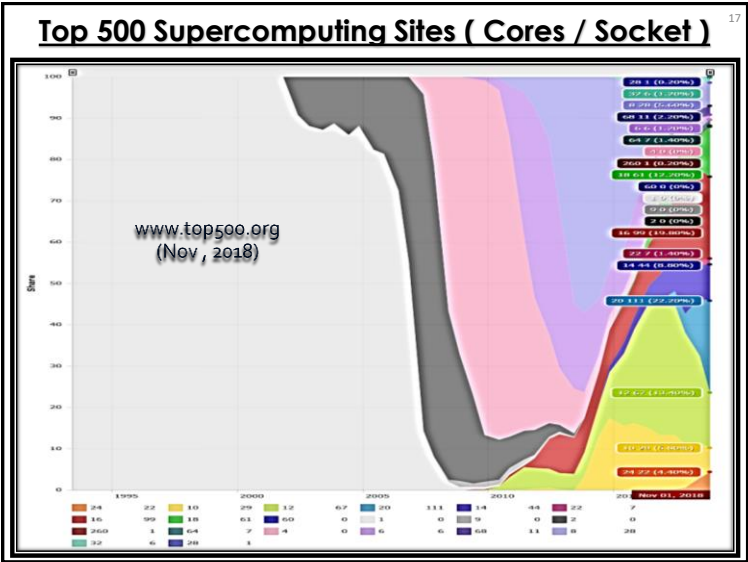


Source: Andrew A. Chien, Vice President of Research, Intel Corporation

Moore's Law Reinterpreted



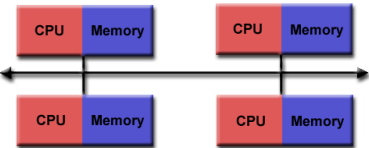
Source: Report of the 2011 Workshop on Exascale Programming Challenges



Some Useful Classifications of Parallel Computers

Parallel Computer Memory Architecture (Distributed Memory)

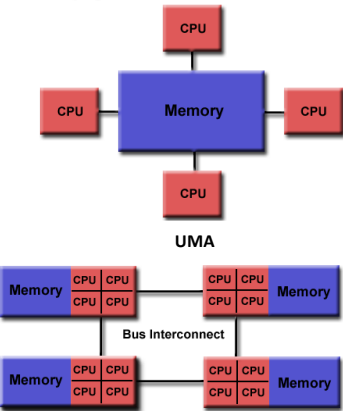
- Each processor has its own local memory — no global address space
- Changes in local memory by one processor have no effect on memory of other processors
- Communication network to connect inter-processor memory
- Programming
 - Message Passing Interface (MPI)
 - Many once available: PVM, Chameleon, MPL, NX, etc.



Source: Blaise Barney, LLNL

Parallel Computer Memory Architecture (Shared Memory)

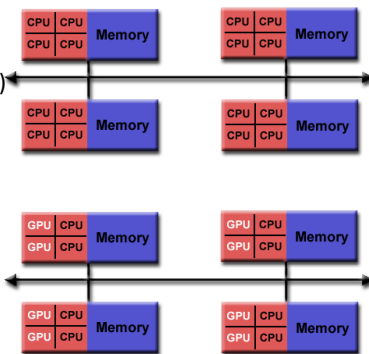
- All processors access all memory as global address space
- Changes in memory by one processor are visible to all others
- Two types
 - Uniform Memory Access (UMA)
 - Non-Uniform Memory Access (NUMA)
- Programming
 - Open Multi-Processing (OpenMP)
 - Cilk/Cilk++ and Intel Cilk Plus
 - Intel Thread Building Block (TBB), etc.



Source: Blaise Barney, LLNL

Parallel Computer Memory Architecture (Hybrid Distributed-Shared Memory)

- The share-memory component can be a cache-coherent SMP or a Graphics Processing Unit (GPU)
- The distributed-memory component is the networking of multiple SMP/GPU machines
- Most common architecture for the largest and fastest computers in the world today
- Programming
 - OpenMP / Cilk + CUDA / OpenCL + MPI, etc.



Source: Blaise Barney, LLNL

Types of Parallelism

$$n C_r = {}^{n-1}C_{r-1} + {}^{n-1}C_r$$

```
int comb ( int n, int r )
{
    if ( r > n ) return 0;
    if ( r == 0 || r == n ) return 1;

    int x, y;

    x = comb( n - 1, r - 1 );
    y = comb( n - 1, r );

    return ( x + y );
}
```

Serial Code

```
comb ( int n, int r )
{
    if ( r > n ) return 0;
    if ( r == 0 || r == n ) return 1;

    int x, y;

    x = spawn comb( n - 1, r - 1 );
    y = comb( n - 1, r );

    sync;

    return ( x + y );
}
```

Parallel Code

Control cannot pass this point until all spawned children have returned.

Grant permission to execute the called (spawned) function in parallel with the caller.

$$\begin{pmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} \end{pmatrix} \xrightarrow[\text{transpose}]{\text{in-place}} \begin{pmatrix} a_{11} & a_{21} & \dots & a_{n1} \\ a_{12} & a_{22} & \dots & a_{n2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \dots & a_{nn} \end{pmatrix}$$

```
for ( int i = 1; i < n; ++i )
    for ( int j = 0; j < i; ++j )
    {
        double t = A[ i ][ j ];
        A[ i ][ j ] = A[ j ][ i ];
        A[ j ][ i ] = t;
    }
```

Serial Code

```
parallel for ( int i = 1; i < n; ++i )
    for ( int j = 0; j < i; ++j )
    {
        double t = A[ i ][ j ];
        A[ i ][ j ] = A[ j ][ i ];
        A[ j ][ i ] = t;
    }
```

Parallel Code

Allows all iterations of the loop to be executed in parallel.

Can be converted to spawns and syncs using recursive divide-and-conquer.

```
parallel for ( int i = s; i < t; ++i )
    BODY( i );
```

divide-and-conquer implementation

```
void recur( int lo, int hi )
{
    if ( hi - lo > GRAINSIZE )
    {
        int mid = lo + ( hi - lo ) / 2;
        spawn recur( lo, mid );
        recur( mid, hi );
        sync;
    }
    else
    {
        for ( int i = lo; i < hi; ++i )
            BODY( i );
    }
}

recur( s, t );
```

Analyzing Parallel Algorithms

Speedup

Let T_p = running time using p identical processing elements

Speedup, $S_p = \frac{T_1}{T_p}$

Theoretically, $S_p \leq p$

Perfect or linear or ideal speedup if $S_p = p$

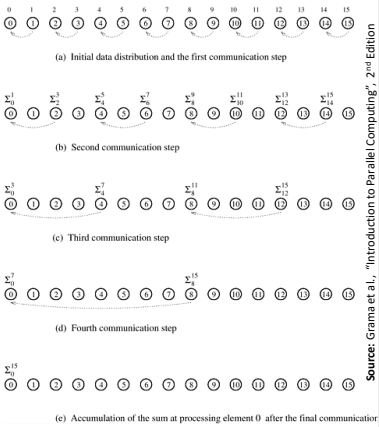
Speedup

Consider adding n numbers using n identical processing elements.

Serial runtime, $T = \Theta(n)$

Parallel runtime, $T_n = \Theta(\log n)$

Speedup, $S_n = \frac{T_1}{T_n} = \Theta\left(\frac{n}{\log n}\right)$



Superlinear Speedup

Theoretically, $S_p \leq p$

But in practice *superlinear speedup* is sometimes observed, that is, $S_p > p$ (why?)

Reasons for superlinear speedup

- Cache effects
- Exploratory decomposition

Parallelism & Span Law

33

We defined, T_p = runtime on p identical processing elements

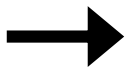
Then span, T_∞ = runtime on an infinite number of identical processing elements

Parallelism, $P = \frac{T_1}{T_\infty}$

Parallelism is an upper bound on speedup, i.e., $S_p \leq P$

Span Law

$T_p \geq T_\infty$



Work Law

34

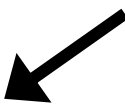
The cost of solving (or work performed for solving) a problem:

On a Serial Computer: is given by T_1

On a Parallel Computer: is given by pT_p

Work Law

$T_p \geq \frac{T_1}{p}$



Bounding Parallel Running Time (T_p)

35

A *runtime/online scheduler* maps tasks to processing elements dynamically at runtime.

A *greedy scheduler* never leaves a processing element idle if it can map a task to it.

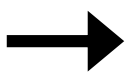
Theorem [Graham'68, Brent'74]: For any greedy scheduler,

$T_p \leq \frac{T_1}{p} + T_\infty$

Corollary: For any greedy scheduler,

$T_p \leq 2T_p^*$

where T_p^* is the running time due to optimal scheduling on p processing elements.



Work Optimality

36

Let T_s = runtime of the optimal or the fastest known serial algorithm

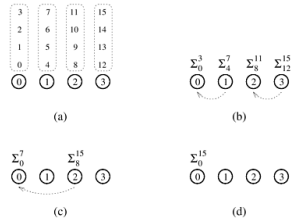
A parallel algorithm is *cost-optimal* or *work-optimal* provided

$pT_p = \Theta(T_s)$

Our algorithm for adding n numbers using n identical processing elements is clearly not work optimal.

Adding n Numbers Work-Optimality

We reduce the number of processing elements which in turn increases the granularity of the subproblem assigned to each processing element.



Suppose we use p processing elements.

First each processing element locally

adds its $\frac{n}{p}$ numbers in time $\Theta\left(\frac{n}{p}\right)$.

Then p processing elements adds these p partial sums in time $\Theta(\log p)$.

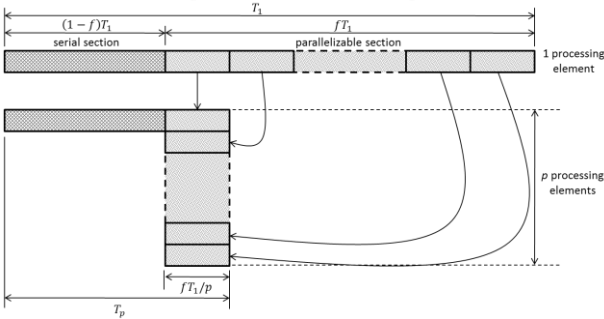
Thus $T_p = \Theta\left(\frac{n}{p} + \log p\right)$, and $T_s = \Theta(n)$.

So the algorithm is work-optimal provided $n = \Omega(p \log p)$.

Source: Grama et al., "Introduction to Parallel Computing", 2nd Edition

Scaling Law

Scaling of Parallel Algorithms (Amdahl's Law)



Suppose only a fraction f of a computation can be parallelized.

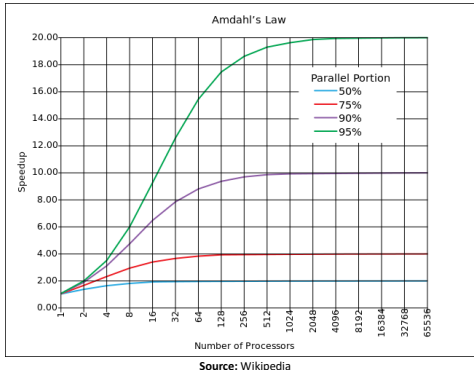
Then parallel running time, $T_p \geq (1 - f)T_1 + f \frac{T_1}{p}$

Speedup, $S_p = \frac{T_1}{T_p} \leq \frac{p}{f + (1-f)p} = \frac{1}{(1-f) + \frac{f}{p}} \leq \frac{1}{1-f}$

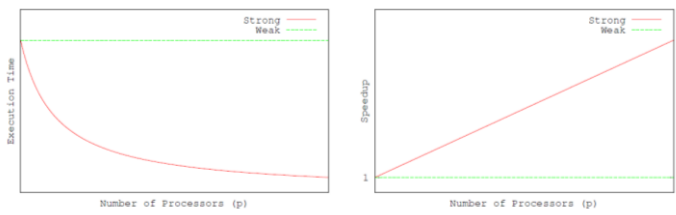
Scaling of Parallel Algorithms (Amdahl's Law)

Suppose only a fraction f of a computation can be parallelized.

Speedup, $S_p = \frac{T_1}{T_p} \leq \frac{1}{(1-f) + \frac{f}{p}} \leq \frac{1}{1-f}$



Strong Scaling vs. Weak Scaling



Strong Scaling

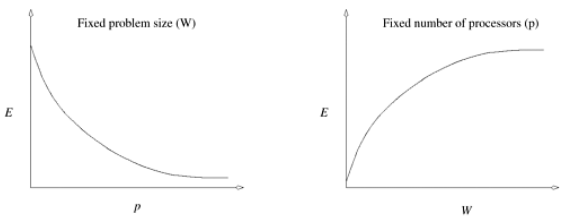
How T_p (or S_p) varies with p when the problem size is fixed.

Weak Scaling

How T_p (or S_p) varies with p when the problem size per processing element is fixed.

Scalable Parallel Algorithms

Efficiency, $E_p = \frac{S_p}{p} = \frac{T_1}{pT_p}$

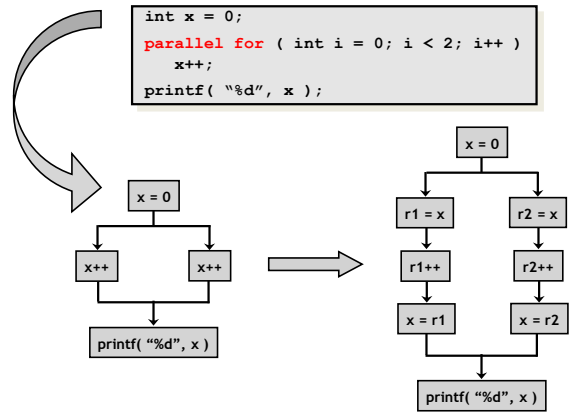


A parallel algorithm is called *scalable* if its efficiency can be maintained at a fixed value by simultaneously increasing the number of processing elements and the problem size. Scalability reflects a parallel algorithm’s ability to utilize increasing processing elements effectively.

Races

Race Bugs

A *determinacy race* occurs if two logically parallel instructions access the same memory location and at least one of them performs a write.



Critical Sections and Mutexes

```
int r = 0;
parallel for ( int i = 0; i < n; i++ )
  r += eval( x[ i ] );
```

race

critical section
two or more strands
must not access
at the same time

```
mutex mtx;
parallel for ( int i = 0; i < n; i++ )
  mtx.lock( );
  r += eval( x[ i ] );
  mtx.unlock( );
```

mutex (mutual exclusion)
an attempt by a strand
to lock an already locked mutex
causes that strand to block (i.e., wait)
until the mutex is unlocked

Problems

- lock overhead
- lock contention

Critical Sections and Mutexes

```
int r = 0;
parallel for ( int i = 0; i < n; i++ )
  r += eval( x[ i ] );
```

race

```
mutex mtx;
parallel for ( int i = 0; i < n; i++ )
  mtx.lock( );
  r += eval( x[ i ] );
  mtx.unlock( );
```

```
mutex mtx;
parallel for ( int i = 0; i < n; i++ )
  int y = eval( x[ i ] );
  mtx.lock( );
  r += y;
  mtx.unlock( );
```

- slightly better solution
- but lock contention can still destroy parallelism

Recursive D&C Implementation of Loops

Recursive D&C Implementation of Parallel Loops

```
parallel for ( int i = s; i < t; ++i )
  BODY( i );
```

divide-and-conquer
implementation

```
void recur( int lo, int hi )
{
  if ( hi - lo > GRAINSIZE )
  {
    int mid = lo + ( hi - lo ) / 2;
    spawn recur( lo, mid );
    recur( mid, hi );
    sync;
  }
  else
  {
    for ( int i = lo; i < hi; ++i )
      BODY( i );
  }
}

recur( s, t );
```

Let $n = t - s$
 m = running time of a single call to BODY
Span: $T_{\infty}(n) = \Theta(\log n + \text{GRAINSIZE} \times m)$

(Lecture 12) Analyzing Parallel Algorithms

Parallel Iterative MM

Iter-MM (Z, X, Y) { X, Y, Z are $n \times n$ matrices,
where n is a positive integer }

- ```

1. for $i \leftarrow 1$ to n do
2. for $j \leftarrow 1$ to n do
3. $Z[i][j] \leftarrow 0$
4. for $k \leftarrow 1$ to n do
5. $Z[i][j] \leftarrow Z[i][j] + X[i][k] \cdot Y[k][j]$

```



**Par-Iter-MM ( Z, X, Y )**      { X, Y, Z are  $n \times n$  matrices,  
where  $n$  is a positive integer }

- ```

1. parallel for  $i \leftarrow 1$  to  $n$  do
2.   parallel for  $j \leftarrow 1$  to  $n$  do
3.      $Z[i][j] \leftarrow 0$ 
4.   for  $k \leftarrow 1$  to  $n$  do
5.      $Z[i][j] \leftarrow Z[i][j] + X[i][k] \cdot Y[k][j]$ 

```

49

Parallel Iterative MM

Par-Iter-MM (Z, X, Y) { X, Y, Z are $n \times n$ matrices,
where n is a positive integer }

- ```

1. parallel for i ← 1 to n do
2. parallel for j ← 1 to n do
3. Z[i][j] ← 0
4. for k ← 1 to n do
5. Z[i][j] ← Z[i][j] + X[i][k] · Y[k][j]

```

**Work:**  $T_1(n) = \Theta(n^3)$

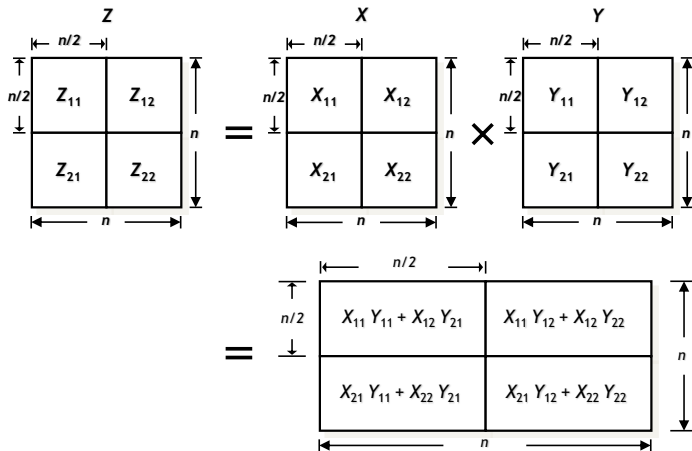
**Span:**  $T_{\infty}(n) = \Theta(n)$

**Parallel Running Time:**  $T_p(n) = O\left(\frac{T_1(n)}{p} + T_\infty(n)\right) = O\left(\frac{n^3}{p} + n\right)$

**Parallelism:**  $\frac{T_1(n)}{T_\infty(n)} = \Theta(n^2)$

50

## Parallel Recursive MM



51

### Parallel Recursive MM

**Par-Rec-MM** ( $Z, X, Y$ ) {  $X, Y, Z$  are  $n \times n$  matrices,  
where  $n = 2^k$  for integer  $k \geq 0$ }

- ```

1. if n = 1 then
2.    $Z \leftarrow Z + X \cdot Y$ 
3. else
4.   spawn Par-Rec-MM (  $Z_{11}, X_{11}, Y_{11}$  )
5.   spawn Par-Rec-MM (  $Z_{12}, X_{11}, Y_{12}$  )
6.   spawn Par-Rec-MM (  $Z_{21}, X_{21}, Y_{11}$  )
7.   Par-Rec-MM (  $Z_{21}, X_{21}, Y_{12}$  )
8.   sync
9.   spawn Par-Rec-MM (  $Z_{11}, X_{12}, Y_{21}$  )
10.  spawn Par-Rec-MM (  $Z_{12}, X_{12}, Y_{22}$  )
11.  spawn Par-Rec-MM (  $Z_{21}, X_{22}, Y_{21}$  )
12.  Par-Rec-MM (  $Z_{22}, X_{22}, Y_{22}$  )
13.  sync
14. endif

```

52

Parallel Recursive MM

53

```
Par-Rec-MM ( Z, X, Y ) { X, Y, Z are n x n matrices,
                        where n = 2^k for integer k ≥ 0 }
1. if n = 1 then
2.   Z ← Z + X · Y
3. else
4.   spawn Par-Rec-MM ( Z11, X11, Y11 )
5.   spawn Par-Rec-MM ( Z12, X11, Y12 )
6.   spawn Par-Rec-MM ( Z21, X21, Y11 )
7.   Par-Rec-MM ( Z21, X21, Y12 )
8.   sync
9.   spawn Par-Rec-MM ( Z11, X12, Y21 )
10.  spawn Par-Rec-MM ( Z12, X12, Y22 )
11.  spawn Par-Rec-MM ( Z21, X22, Y21 )
12.  Par-Rec-MM ( Z22, X22, Y22 )
13.  sync
14. endif
```

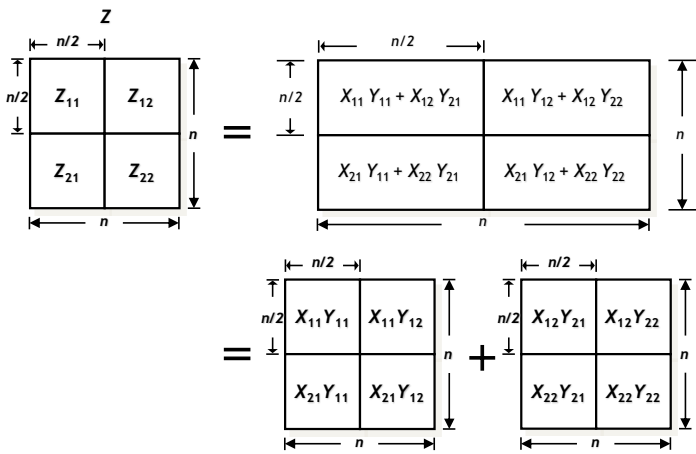
Work:
$$T_1(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ 8T_1\left(\frac{n}{2}\right) + \Theta(1), & \text{otherwise.} \end{cases}$$
$$= \Theta(n^3) \quad [\text{MT Case 1}]$$

Span:
$$T_\infty(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ 2T_\infty\left(\frac{n}{2}\right) + \Theta(1), & \text{otherwise.} \end{cases}$$
$$= \Theta(n) \quad [\text{MT Case 1}]$$

Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta(n^2)$
Additional Space:
 $S_\infty(n) = \Theta(1)$

Recursive MM with More Parallelism

54



Recursive MM with More Parallelism

55

```
Par-Rec-MM2 ( Z, X, Y ) { X, Y, Z are n x n matrices,
                        where n = 2^k for integer k ≥ 0 }
1. if n = 1 then
2.   Z ← Z + X · Y
3. else { T is a temporary n x n matrix }
4.   spawn Par-Rec-MM2 ( Z11, X11, Y11 )
5.   spawn Par-Rec-MM2 ( Z12, X11, Y12 )
6.   spawn Par-Rec-MM2 ( Z21, X21, Y11 )
7.   spawn Par-Rec-MM2 ( Z21, X21, Y12 )
8.   spawn Par-Rec-MM2 ( T11, X12, Y21 )
9.   spawn Par-Rec-MM2 ( T12, X12, Y22 )
10.  spawn Par-Rec-MM2 ( T21, X22, Y21 )
11.  Par-Rec-MM2 ( T22, X22, Y22 )
12.  sync
13.  parallel for i ← 1 to n do
14.    parallel for j ← 1 to n do
15.      Z[i][j] ← Z[i][j] + T[i][j]
16.  endif
```

Recursive MM with More Parallelism

56

```
Par-Rec-MM2 ( Z, X, Y ) { X, Y, Z are n x n matrices,
                        where n = 2^k for integer k ≥ 0 }
1. if n = 1 then
2.   Z ← Z + X · Y
3. else { T is a temporary n x n matrix }
4.   spawn Par-Rec-MM2 ( Z11, X11, Y11 )
5.   spawn Par-Rec-MM2 ( Z12, X11, Y12 )
6.   spawn Par-Rec-MM2 ( Z21, X21, Y11 )
7.   spawn Par-Rec-MM2 ( Z21, X21, Y12 )
8.   spawn Par-Rec-MM2 ( T11, X12, Y21 )
9.   spawn Par-Rec-MM2 ( T12, X12, Y22 )
10.  spawn Par-Rec-MM2 ( T21, X22, Y21 )
11.  Par-Rec-MM2 ( T22, X22, Y22 )
12.  sync
13.  parallel for i ← 1 to n do
14.    parallel for j ← 1 to n do
15.      Z[i][j] ← Z[i][j] + T[i][j]
16.  endif
```

Work:
$$T_1(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ 8T_1\left(\frac{n}{2}\right) + \Theta(n^2), & \text{otherwise.} \end{cases}$$
$$= \Theta(n^3) \quad [\text{MT Case 1}]$$

Span:
$$T_\infty(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ T_\infty\left(\frac{n}{2}\right) + \Theta(\log n), & \text{otherwise.} \end{cases}$$
$$= \Theta(\log^2 n) \quad [\text{MT Case 2}]$$

Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta\left(\frac{n^3}{\log^2 n}\right)$
Additional Space:
$$S_\infty(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ 8S_\infty\left(\frac{n}{2}\right) + \Theta(n^2), & \text{otherwise.} \end{cases}$$
$$= \Theta(n^3) \quad [\text{MT Case 1}]$$

57

Parallel Merge Sort

58

Parallel Merge Sort

Merge-Sort (A, p, r) { sort the elements in A[p ... r] }

1. if $p < r$ then

2. $q \leftarrow \lfloor (p + r) / 2 \rfloor$

3. Merge-Sort (A, p, q)

4. Merge-Sort (A, q + 1, r)

5. Merge (A, p, q, r)

Par-Merge-Sort (A, p, r) { sort the elements in A[p ... r] }

1. if $p < r$ then

2. $q \leftarrow \lfloor (p + r) / 2 \rfloor$

3. spawn Merge-Sort (A, p, q)

4. Merge-Sort (A, q + 1, r)

5. sync

6. Merge (A, p, q, r)

59

Parallel Merge Sort

Par-Merge-Sort (A, p, r) { sort the elements in A[p ... r] }

1. if $p < r$ then

2. $q \leftarrow \lfloor (p + r) / 2 \rfloor$

3. spawn Merge-Sort (A, p, q)

4. Merge-Sort (A, q + 1, r)

5. sync

6. Merge (A, p, q, r)

Work: $T_1(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ 2T_1(\frac{n}{2}) + \Theta(n), & \text{otherwise.} \end{cases}$

$= \Theta(n \log n)$ [MT Case 2]

Span: $T_\infty(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ T_\infty(\frac{n}{2}) + \Theta(n), & \text{otherwise.} \end{cases}$

$= \Theta(n)$ [MT Case 3]

Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta(\log n)$

60

Parallel Merge

$n_1 = r_1 - p_1 + 1$

$n_2 = r_2 - p_2 + 1$

subarrays to merge: $T[p_1..r_1]$ $T[p_2..r_2]$

$\leq x$ $\geq x$

$< x$ $\geq x$

suppose: $n_1 \geq n_2$

merge

copy

merge

merged output: $A[p_3..r_3]$
 $n_3 = r_3 - p_3 + 1 = n_1 + n_2$

Source: Cormen et al., "Introduction to Algorithms", 3rd Edition

15

Parallel Merge

$n_1 = r_1 - p_1 + 1$ $n_2 = r_2 - p_2 + 1$

subarrays to merge: $T[p_1..r_1]$ $T[p_2..r_2]$

T ... $\leq x$ x $\geq x$... $< x$ $\geq x$...

suppose: $n_1 \geq n_2$

merge copy merge

A ... $\leq x$ x $\geq x$...

merged output: $A[p_3..r_3]$
 $n_3 = r_3 - p_3 + 1 = n_1 + n_2$

Step 1: Find $x = T[q_1]$, where q_1 is the midpoint of $T[p_1..r_1]$

61

Source: Cormen et al., "Introduction to Algorithms", 3rd Edition

Parallel Merge

$n_1 = r_1 - p_1 + 1$ $n_2 = r_2 - p_2 + 1$

subarrays to merge: $T[p_1..r_1]$ $T[p_2..r_2]$

T ... $\leq x$ x $\geq x$... $< x$ $\geq x$...

suppose: $n_1 \geq n_2$

merge copy merge

A ... $\leq x$ x $\geq x$...

merged output: $A[p_3..r_3]$
 $n_3 = r_3 - p_3 + 1 = n_1 + n_2$

Step 2: Use binary search to find the index q_2 in subarray $T[p_2..r_2]$ so that the subarray would still be sorted if we insert x between $T[q_2 - 1]$ and $T[q_2]$

62

Source: Cormen et al., "Introduction to Algorithms", 3rd Edition

Parallel Merge

$n_1 = r_1 - p_1 + 1$ $n_2 = r_2 - p_2 + 1$

subarrays to merge: $T[p_1..r_1]$ $T[p_2..r_2]$

T ... $\leq x$ x $\geq x$... $< x$ $\geq x$...

suppose: $n_1 \geq n_2$

merge copy merge

A ... $\leq x$ x $\geq x$...

merged output: $A[p_3..r_3]$
 $n_3 = r_3 - p_3 + 1 = n_1 + n_2$

Step 3: Copy x to $A[q_3]$, where $q_3 = p_3 + (q_1 - p_1) + (q_2 - p_2)$

63

Source: Cormen et al., "Introduction to Algorithms", 3rd Edition

Parallel Merge

$n_1 = r_1 - p_1 + 1$ $n_2 = r_2 - p_2 + 1$

subarrays to merge: $T[p_1..r_1]$ $T[p_2..r_2]$

T ... $\leq x$ x $\geq x$... $< x$ $\geq x$...

suppose: $n_1 \geq n_2$

merge copy merge

A ... $\leq x$ x $\geq x$...

merged output: $A[p_3..r_3]$
 $n_3 = r_3 - p_3 + 1 = n_1 + n_2$

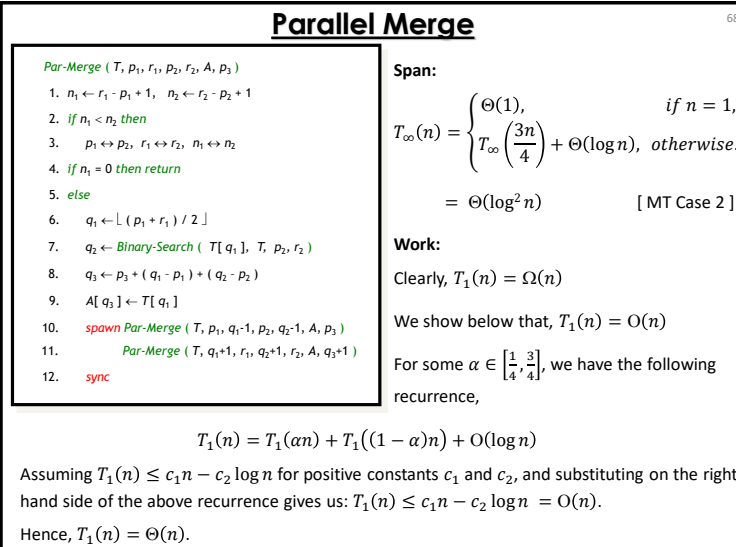
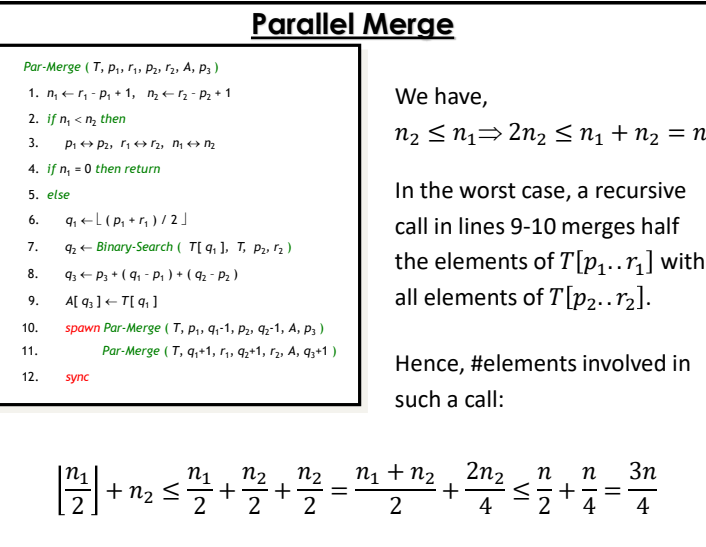
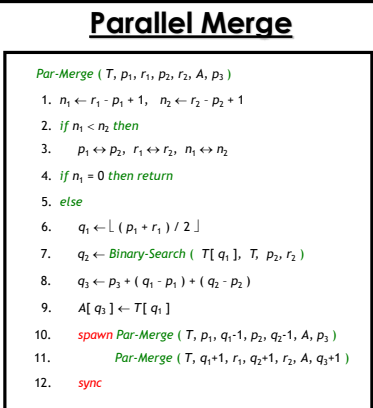
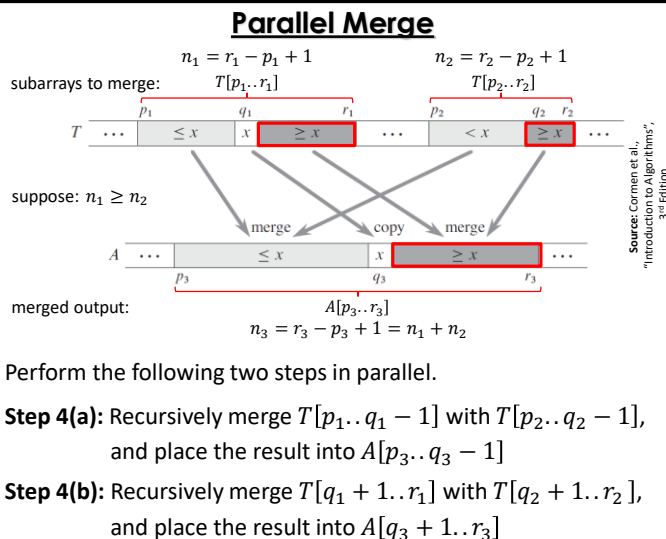
Perform the following two steps in parallel.

Step 4(a): Recursively merge $T[p_1..q_1 - 1]$ with $T[p_2..q_2 - 1]$, and place the result into $A[p_3..q_3 - 1]$

64

Source: Cormen et al., "Introduction to Algorithms", 3rd Edition

(Lecture 12) Analyzing Parallel Algorithms



Parallel Merge Sort with Parallel Merge

```
Par-Merge-Sort ( A, p, r ) { sort the elements in A[ p ... r ] }  
1. if p < r then  
2.   q ← ⌊ ( p + r ) / 2 ⌋  
3.   spawn Merge-Sort ( A, p, q )  
4.   Merge-Sort ( A, q + 1, r )  
5.   sync  
6.   Par-Merge ( A, p, q, r )
```

Work: $T_1(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ 2T_1\left(\frac{n}{2}\right) + \Theta(n), & \text{otherwise.} \end{cases}$
 $= \Theta(n \log n)$ [MT Case 2]
Span: $T_\infty(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ T_\infty\left(\frac{n}{2}\right) + \Theta(\log^2 n), & \text{otherwise.} \end{cases}$
 $= \Theta(\log^3 n)$ [MT Case 2]
Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta\left(\frac{n}{\log^2 n}\right)$

Parallel Prefix Sums

Parallel Prefix Sums

Input: A sequence of n elements $\{x_1, x_2, \dots, x_n\}$ drawn from a set S with a binary associative operation, denoted by \oplus .
Output: A sequence of n partial sums $\{s_1, s_2, \dots, s_n\}$, where $s_i = x_1 \oplus x_2 \oplus \dots \oplus x_i$ for $1 \leq i \leq n$.

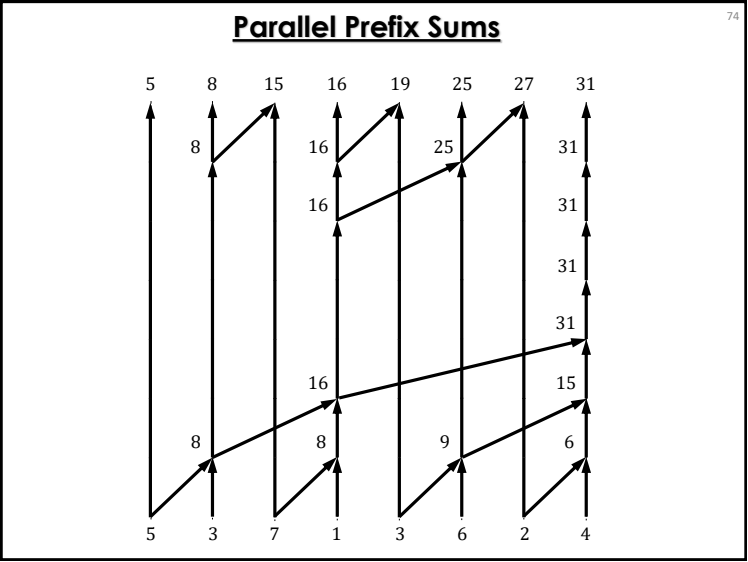
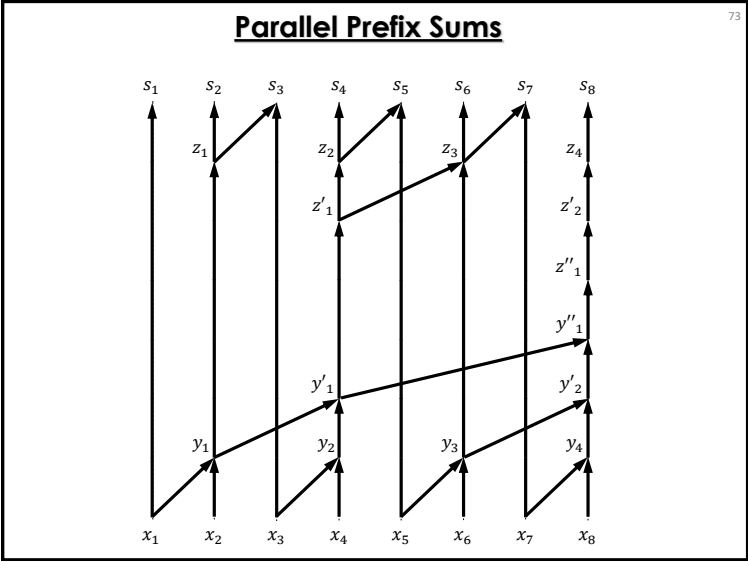
x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8
5	3	7	1	3	6	2	4

\oplus = binary addition

5	8	15	16	19	25	27	31
s_1	s_2	s_3	s_4	s_5	s_6	s_7	s_8

Parallel Prefix Sums

```
Prefix-Sum (  $\langle x_1, x_2, \dots, x_n \rangle, \oplus$  ) {  $n = 2^k$  for some  $k \geq 0$ .  
Return prefix sums  
 $\langle s_1, s_2, \dots, s_n \rangle$  }  
1. if  $n = 1$  then  
2.    $s_1 \leftarrow x_1$   
3. else  
4.   parallel for  $i \leftarrow 1$  to  $n/2$  do  
5.      $y_i \leftarrow x_{2i-1} \oplus x_{2i}$   
6.    $\langle z_1, z_2, \dots, z_{n/2} \rangle \leftarrow \text{Prefix-Sum}(\langle y_1, y_2, \dots, y_{n/2} \rangle, \oplus)$   
7.   parallel for  $i \leftarrow 1$  to  $n$  do  
8.     if  $i = 1$  then  $s_1 \leftarrow x_1$   
9.     else if  $i$  is even then  $s_i \leftarrow z_{i/2}$   
10.    else  $s_i \leftarrow z_{(i-1)/2} \oplus x_i$   
11. return  $\langle s_1, s_2, \dots, s_n \rangle$ 
```



75

Parallel Prefix Sums

```
Prefix-Sum (  $\langle x_1, x_2, \dots, x_n \rangle, \oplus$  )  [  $n = 2^k$  for some  $k \geq 0$ .  
  Return prefix sums  
   $\langle s_1, s_2, \dots, s_n \rangle$  ]  
1. if  $n = 1$  then  
2.    $s_1 \leftarrow x_1$   
3. else  
4.   parallel for  $i \leftarrow 1$  to  $n/2$  do  
5.      $y_i \leftarrow x_{2i-1} \oplus x_{2i}$   
6.    $\langle z_1, z_2, \dots, z_{n/2} \rangle \leftarrow \text{Prefix-Sum}(\langle y_1, y_2, \dots, y_{n/2} \rangle, \oplus)$   
7.   parallel for  $i \leftarrow 1$  to  $n$  do  
8.     if  $i = 1$  then  $s_i \leftarrow x_1$   
9.     else if  $i$  is even then  $s_i \leftarrow z_{i/2}$   
10.    else  $s_i \leftarrow z_{(i-1)/2} \oplus x_i$   
11. return  $\langle s_1, s_2, \dots, s_n \rangle$ 
```

Work:

$$T_1(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ T_1\left(\frac{n}{2}\right) + \Theta(n), & \text{otherwise.} \end{cases}$$
$$= \Theta(n)$$

Span:

$$T_\infty(n) = \begin{cases} \Theta(1), & \text{if } n = 1, \\ T_\infty\left(\frac{n}{2}\right) + \Theta(1), & \text{otherwise.} \end{cases}$$
$$= \Theta(\log n)$$

Parallelism:

$$\frac{T_1(n)}{T_\infty(n)} = \Theta\left(\frac{n}{\log n}\right)$$

Observe that we have assumed here that a *parallel for* loop can be executed in $\Theta(1)$ time. But recall that *cilk_for* is implemented using divide-and-conquer, and so in practice, it will take $\Theta(\log n)$ time. In that case, we will have $T_\infty(n) = \Theta(\log^2 n)$, and parallelism = $\Theta(n/\log^2 n)$.

76

Parallel Partition

Parallel Partition77

Input: An array $A[q : r]$ of distinct elements, and an element x from $A[q : r]$.

Output: Rearrange the elements of $A[q : r]$, and return an index $k \in [q, r]$, such that all elements in $A[q : k - 1]$ are smaller than x , all elements in $A[k + 1 : r]$ are larger than x , and $A[k] = x$.

```
Par-Partition ( A[ q : r ], x )
1. n ← r − q + 1
2. if n = 1 then return q
3. array B[ 0: n − 1 ], lt[ 0: n − 1 ], gt[ 0: n − 1 ]
4. parallel for i ← 0 to n − 1 do
5.   B[ i ] ← A[ q + i ]
6.   if B[ i ] < x then lt[ i ] ← 1 else lt[ i ] ← 0
7.   if B[ i ] > x then gt[ i ] ← 1 else gt[ i ] ← 0
8. lt[ 0: n − 1 ] ← Par-Prefix-Sum ( lt[ 0: n − 1 ], + )
9. gt[ 0: n − 1 ] ← Par-Prefix-Sum ( gt[ 0: n − 1 ], + )
10. k ← q + lt[ n − 1 ], A[ k ] ← x
11. parallel for i ← 0 to n − 1 do
12.   if B[ i ] < x then A[ q + lt[ i ] − 1 ] ← B[ i ]
13.   else if B[ i ] > x then A[ k + gt[ i ] ] ← B[ i ]
14. return k
```

Parallel Partition78

A:

9	5	7	11	1	3	8	14	4	21
---	---	---	----	---	---	---	----	---	----

 x = 8

Parallel Partition79

A:

9	5	7	11	1	3	8	14	4	21
---	---	---	----	---	---	---	----	---	----

 x = 8

B:

9	5	7	11	1	3	8	14	4	21
---	---	---	----	---	---	---	----	---	----

lt:

0	1	1	0	1	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---

 gt:

1	0	0	1	0	0	0	1	0	1
---	---	---	---	---	---	---	---	---	---

Parallel Partition80

A:

9	5	7	11	1	3	8	14	4	21
---	---	---	----	---	---	---	----	---	----

 x = 8

B:

9	5	7	11	1	3	8	14	4	21
---	---	---	----	---	---	---	----	---	----

lt:

0	1	1	0	1	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---

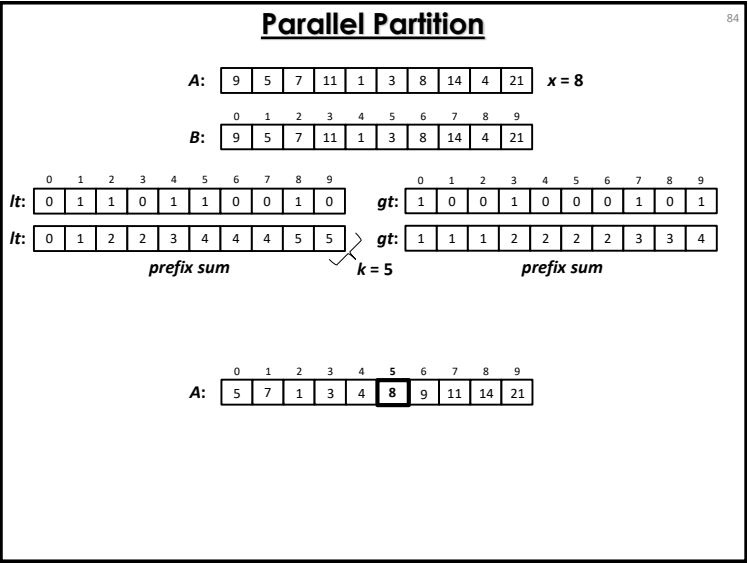
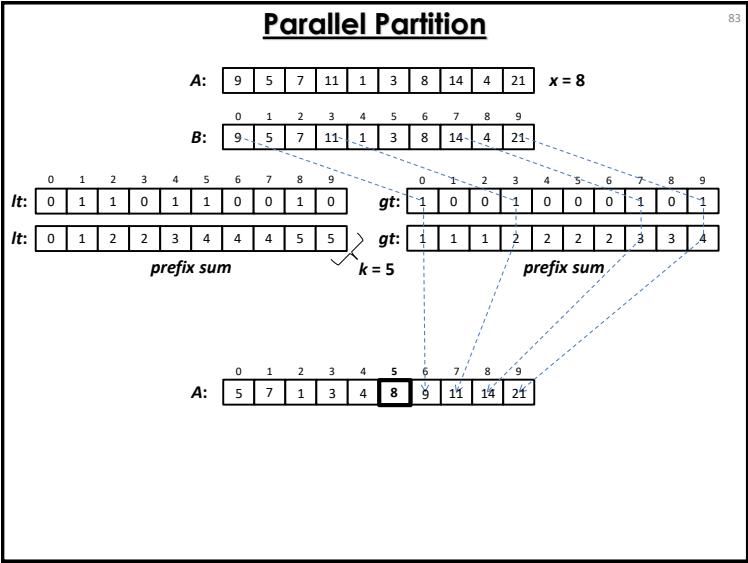
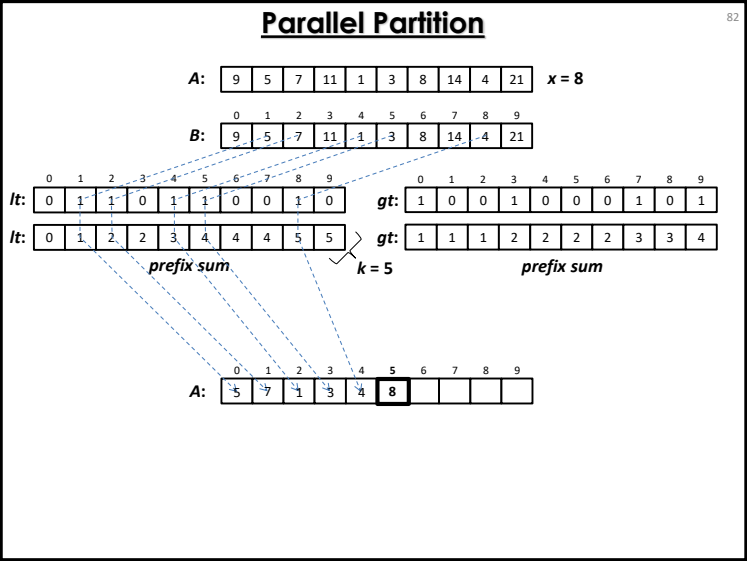
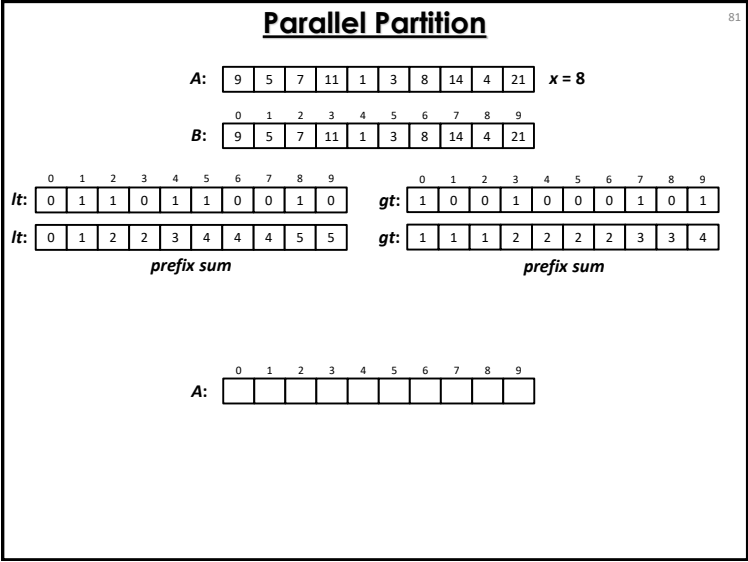
 prefix sum

gt:

1	0	0	1	0	0	0	1	0	1
---	---	---	---	---	---	---	---	---	---

 prefix sum

(Lecture 12) Analyzing Parallel Algorithms



Parallel Partition: Analysis

85

```
Par-Partition ( A[ q : r ], x )
1. n ← r - q + 1
2. if n = 1 then return q
3. array B[ 0 : n - 1 ], lt[ 0 : n - 1 ], gt[ 0 : n - 1 ]
4. parallel for i ← 0 to n - 1 do
5.   B[ i ] ← A[ q + i ]
6.   if B[ i ] < x then lt[ i ] ← 1 else lt[ i ] ← 0
7.   if B[ i ] > x then gt[ i ] ← 1 else gt[ i ] ← 0
8. lt[ 0 : n - 1 ] ← Par-Prefix-Sum ( lt[ 0 : n - 1 ], + )
9. gt[ 0 : n - 1 ] ← Par-Prefix-Sum ( gt[ 0 : n - 1 ], + )
10. k ← q + lt[ n - 1 ], A[ k ] ← x
11. parallel for i ← 0 to n - 1 do
12.   if B[ i ] < x then A[ q + lt[ i ] - 1 ] ← B[ i ]
13.   else if B[ i ] > x then A[ k + gt[ i ] ] ← B[ i ]
14. return k
```

Work:
 $T_1(n) = \Theta(n)$ [lines 1 – 7]
 $\quad + \Theta(n)$ [lines 8 – 9]
 $\quad + \Theta(n)$ [lines 10 – 14]
 $\quad = \Theta(n)$

Span:
Assuming $\log n$ depth for *parallel* for loops:
 $T_\infty(n) = \Theta(\log n)$ [lines 1 – 7]
 $\quad + \Theta(\log^2 n)$ [lines 8 – 9]
 $\quad + \Theta(\log n)$ [lines 10 – 14]
 $\quad = \Theta(\log^2 n)$

Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta\left(\frac{n}{\log^2 n}\right)$



Parallel Quicksort

86

Randomized Parallel QuickSort

87

Input: An array $A[q : r]$ of distinct elements.
Output: Elements of $A[q : r]$ sorted in increasing order of value.

```
Par-Randomized-QuickSort ( A[ q : r ] )
1. n ← r - q + 1
2. if n ≤ 30 then
3.   sort A[ q : r ] using any sorting algorithm
4. else
5.   select a random element x from A[ q : r ]
6.   k ← Par-Partition ( A[ q : r ], x )
7.   spawn Par-Randomized-QuickSort ( A[ q : k - 1 ] )
8.   Par-Randomized-QuickSort ( A[ k + 1 : r ] )
9. sync
```



Randomized Parallel QuickSort: Analysis

88

```
Par-Randomized-QuickSort ( A[ q : r ] )
1. n ← r - q + 1
2. if n ≤ 30 then
3.   sort A[ q : r ] using any sorting algorithm
4. else
5.   select a random element x from A[ q : r ]
6.   k ← Par-Partition ( A[ q : r ], x )
7.   spawn Par-Randomized-QuickSort ( A[ q : k - 1 ] )
8.   Par-Randomized-QuickSort ( A[ k + 1 : r ] )
9. sync
```

Lines 1—6 take $\Theta(\log^2 n)$ parallel time and perform $\Theta(n)$ work.
Also the recursive spawns in lines 7—8 work on disjoint parts of $A[q : r]$. So the upper bounds on the parallel time and the total work in each level of recursion are $\Theta(\log^2 n)$ and $\Theta(n)$, respectively.
Hence, if D is the *recursion depth* of the algorithm, then
 $T_1(n) = O(nD)$ and $T_\infty(n) = O(D \log^2 n)$

Randomized Parallel QuickSort: Analysis

```
Par-Randomized-QuickSort ( A[ q : r ] )
1.  n ← r − q + 1
2.  if n ≤ 30 then
3.    sort A[ q : r ] using any sorting algorithm
4.  else
5.    select a random element x from A[ q : r ]
6.    k ← Par-Partition ( A[ q : r ], x )
7.    spawn Par-Randomized-QuickSort ( A[ q : k − 1 ] )
8.    Par-Randomized-QuickSort ( A[ k + 1 : r ] )
9.    sync
```

We already proved that w.h.p. recursion depth, $D = O(\log n)$.
Hence, with high probability,
 $T_1(n) = O(n \log n)$ and $T_\infty(n) = O(\log^3 n)$

